

INRAe



université  
PARIS-SACLAY

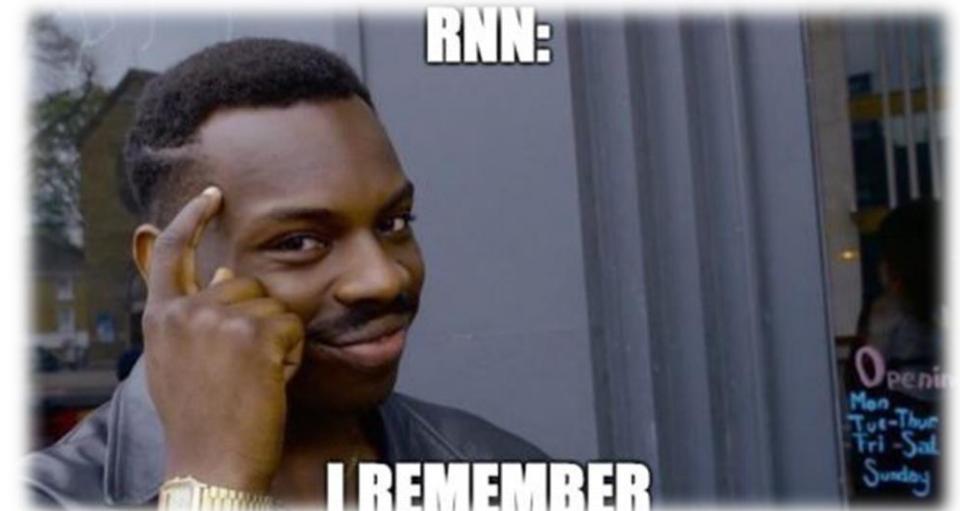
# → Recurrent Neural Networks

Alberto TONDA, Ph.D. (Senior permanent researcher, DR)

UMR 518 MIA-PS, INRAE, AgroParisTech, Université Paris-Saclay  
UAR 3611, Institut des Systèmes Complexes de Paris Île-de-France

# > Outline

- Dynamical systems
- What is a Recurrent Neural Network (RNN)?
- First architectures and issues
- Long-Short-Term Memory Networks (LSTM)



# > Dynamical systems

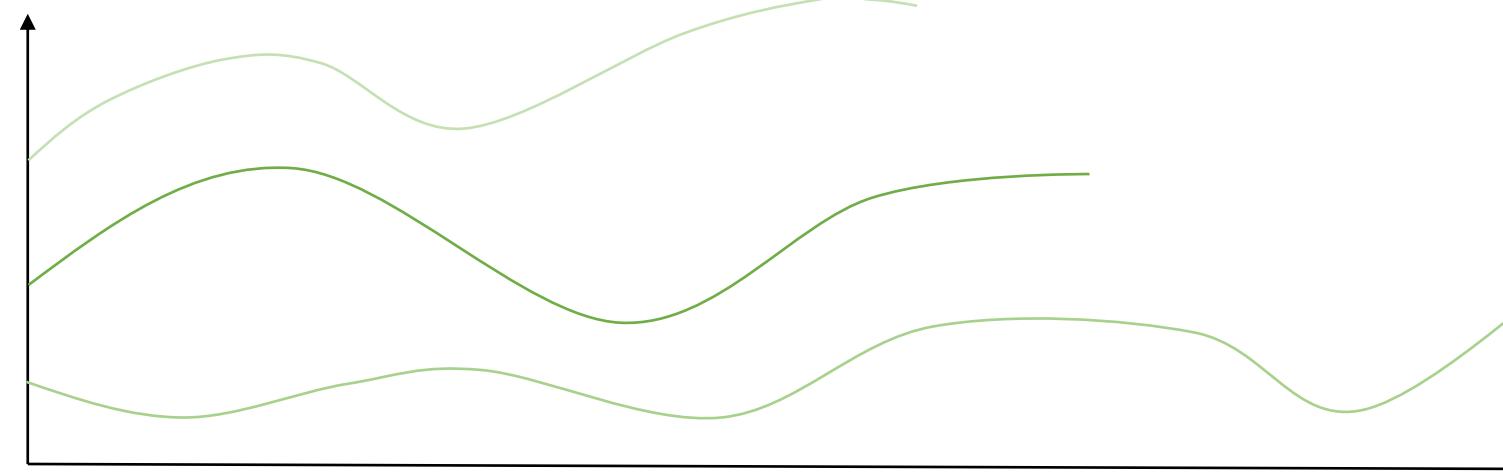
- Systems where the output does not depend only on input
  - But also on a history of *previous inputs*
  - Typical of **time-series** analysis and forecasting
  - But in general, any type of sequence
- $x_{t+1} = f(x_t, h_t)$ , with  $h_t = g(x_0, x_1, \dots, x_{t-1})$
- $h_t$  can also be called **state**, **hidden state**, history of system

# > Dynamical systems

- Alternatives to computing and storing  $h_t$ ?
- Autoregressive models
  - Assumption:  $x_{t+1} = f(x_t, h_t)$ , but  $h_t = g(x_{t-1}, x_{t-2}, \dots, x_{t-N})$
  - In other words,  $x_{t+1} = f(x_t, x_{t-1}, x_{t-2}, \dots, x_{t-N})$
- They might look simple, but good performance
  - Auto-Regressive Moving Average (ARMA)
  - Auto-Regressive Integrated Moving Average (ARIMA)

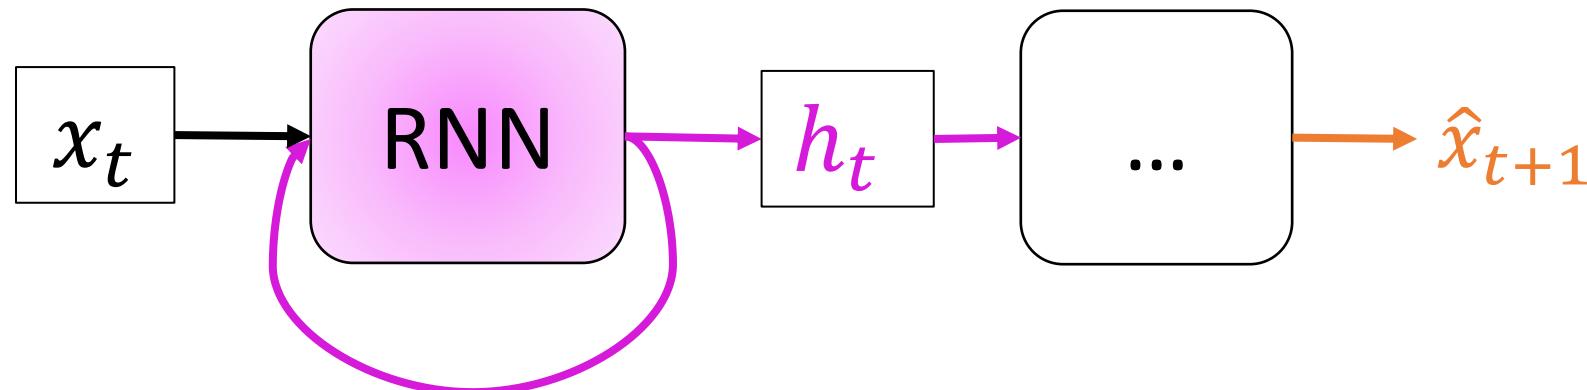
# > Dynamical systems

- For neural networks, this could be an issue
  - So far, input tensors for an application had the **exact same shape**
  - *Variable number of steps* in sequences for same application
  - For example, how could a CNN deal with a variable-sized input?



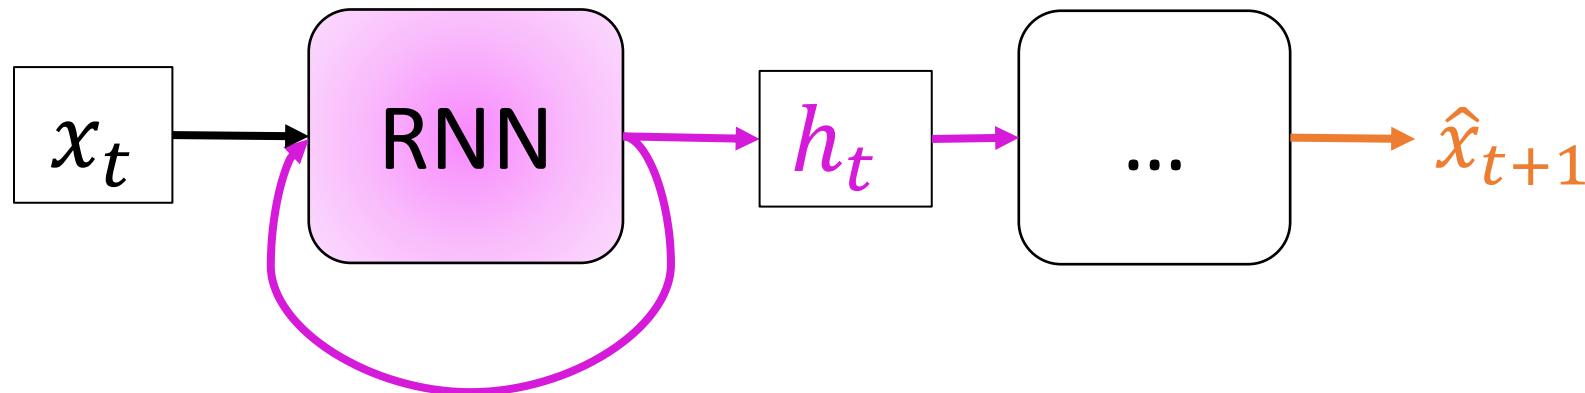
# > Recurrent Neural Networks

- Interestingly, first ideas are from 1925 (!)
  - “Ising model” (Lenz and Ising), model of magnetism (no learning)
  - Shin’ici Amari, 1972, version with adaptable weights
  - Popularized by John Hopfield in 1982, “Hopfield Networks”



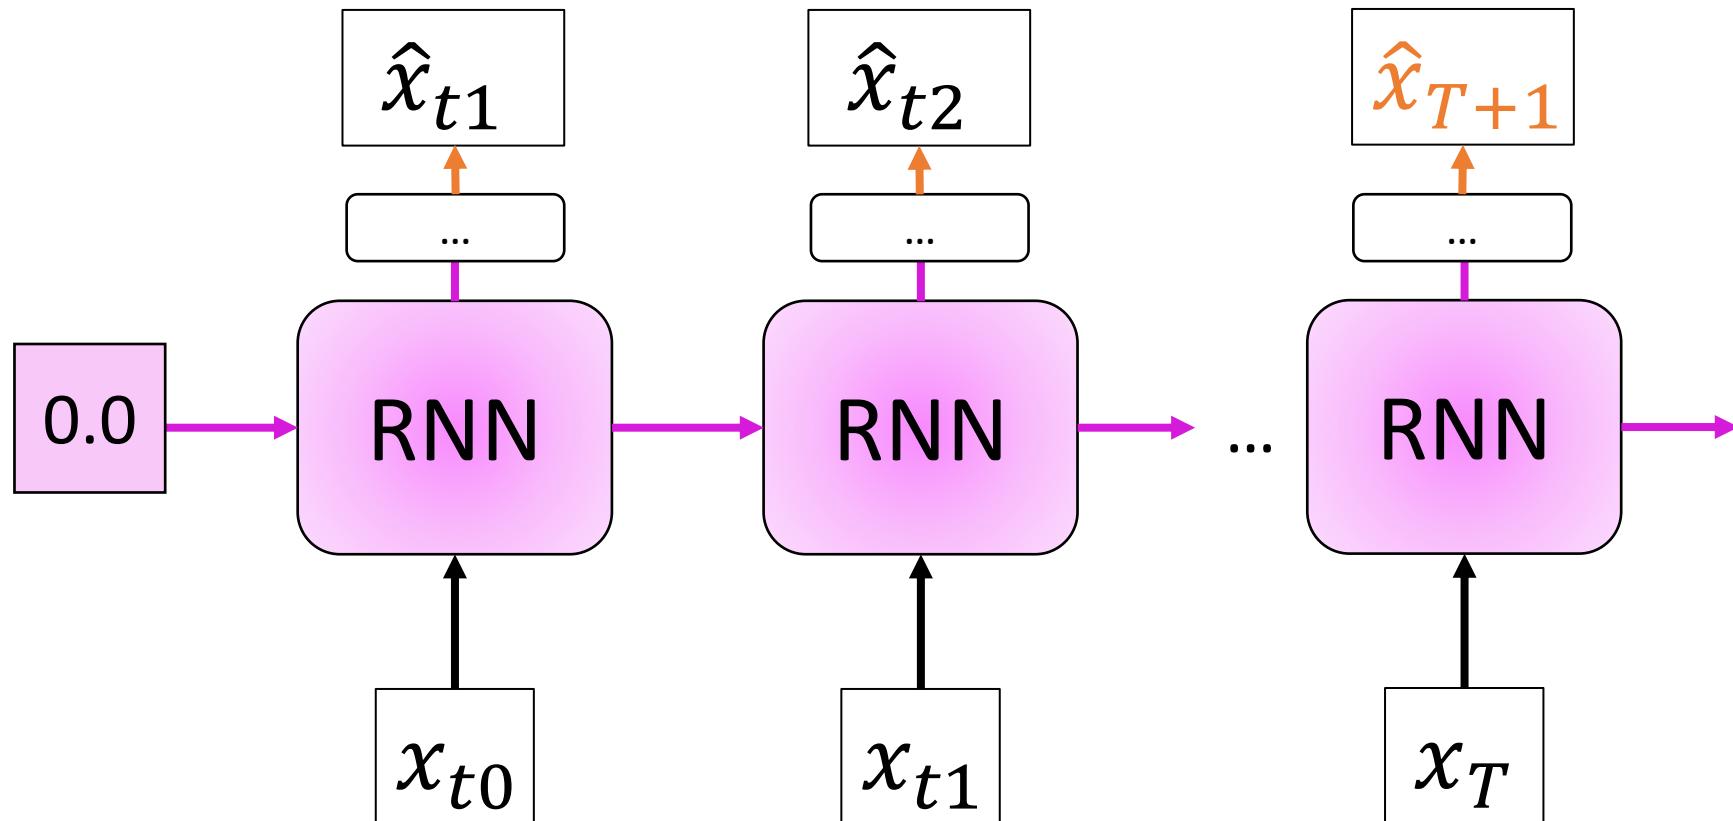
# > Recurrent Neural Networks

- The architecture is conceived to predict  $x_{t+1} = f(x_t, h_t)$ 
  - RNN unit is designed to compute  $h_t = f(x_t, h_{t-1})$
  - Other modules can later obtain  $x_{t+1} = f(x_t, h_t)$



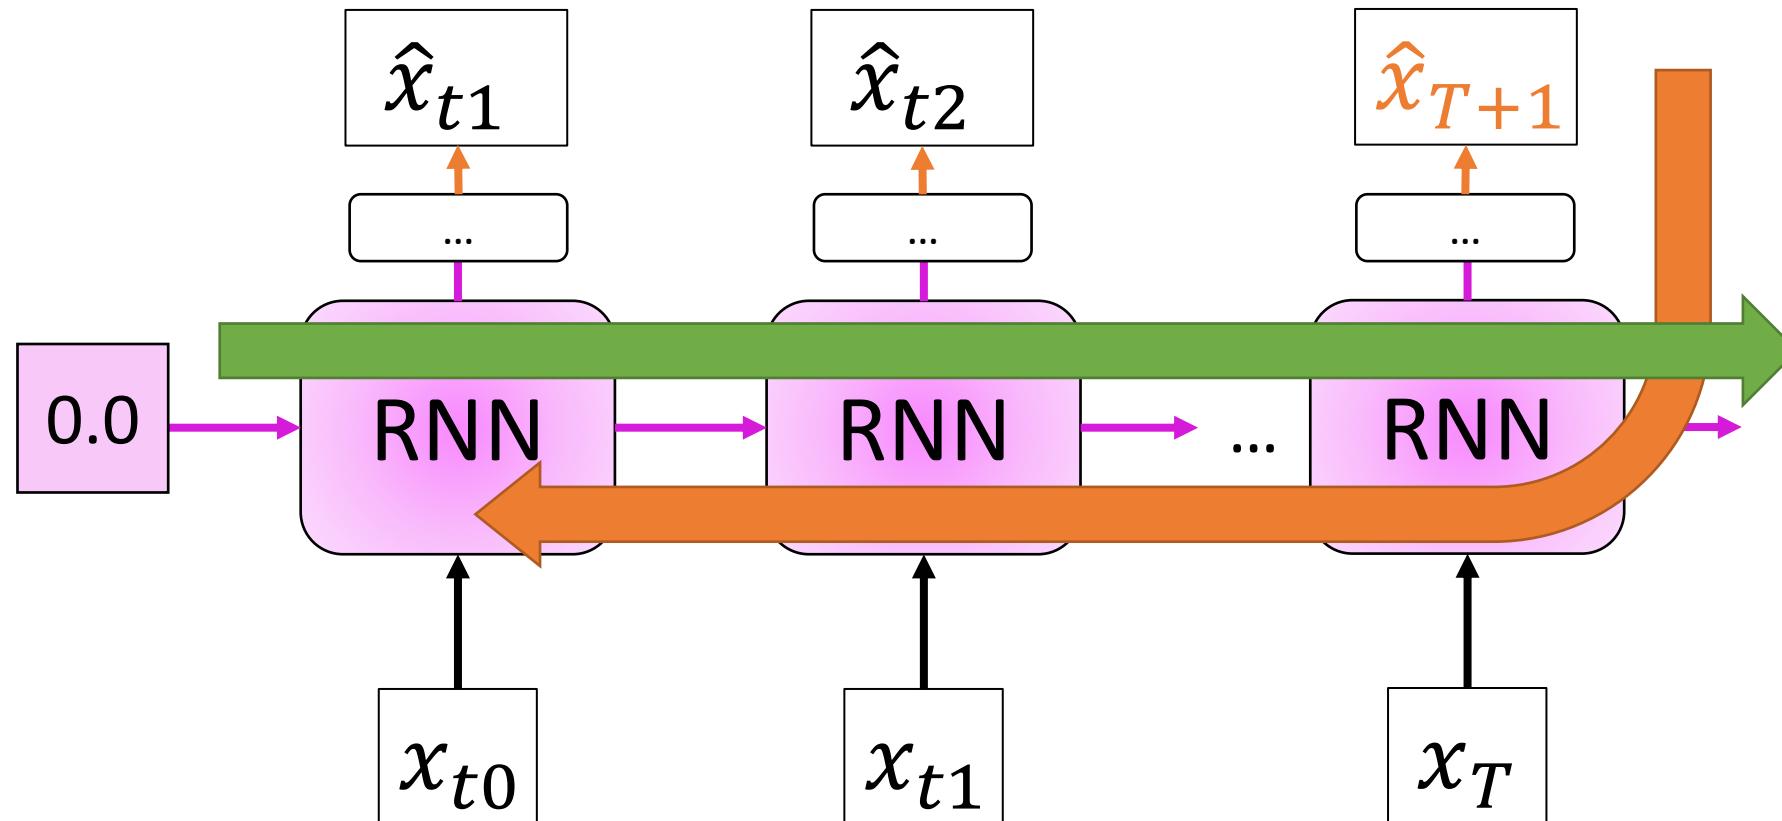
# > Recurrent Neural Networks

- How is a RNN trained? Unrolling!



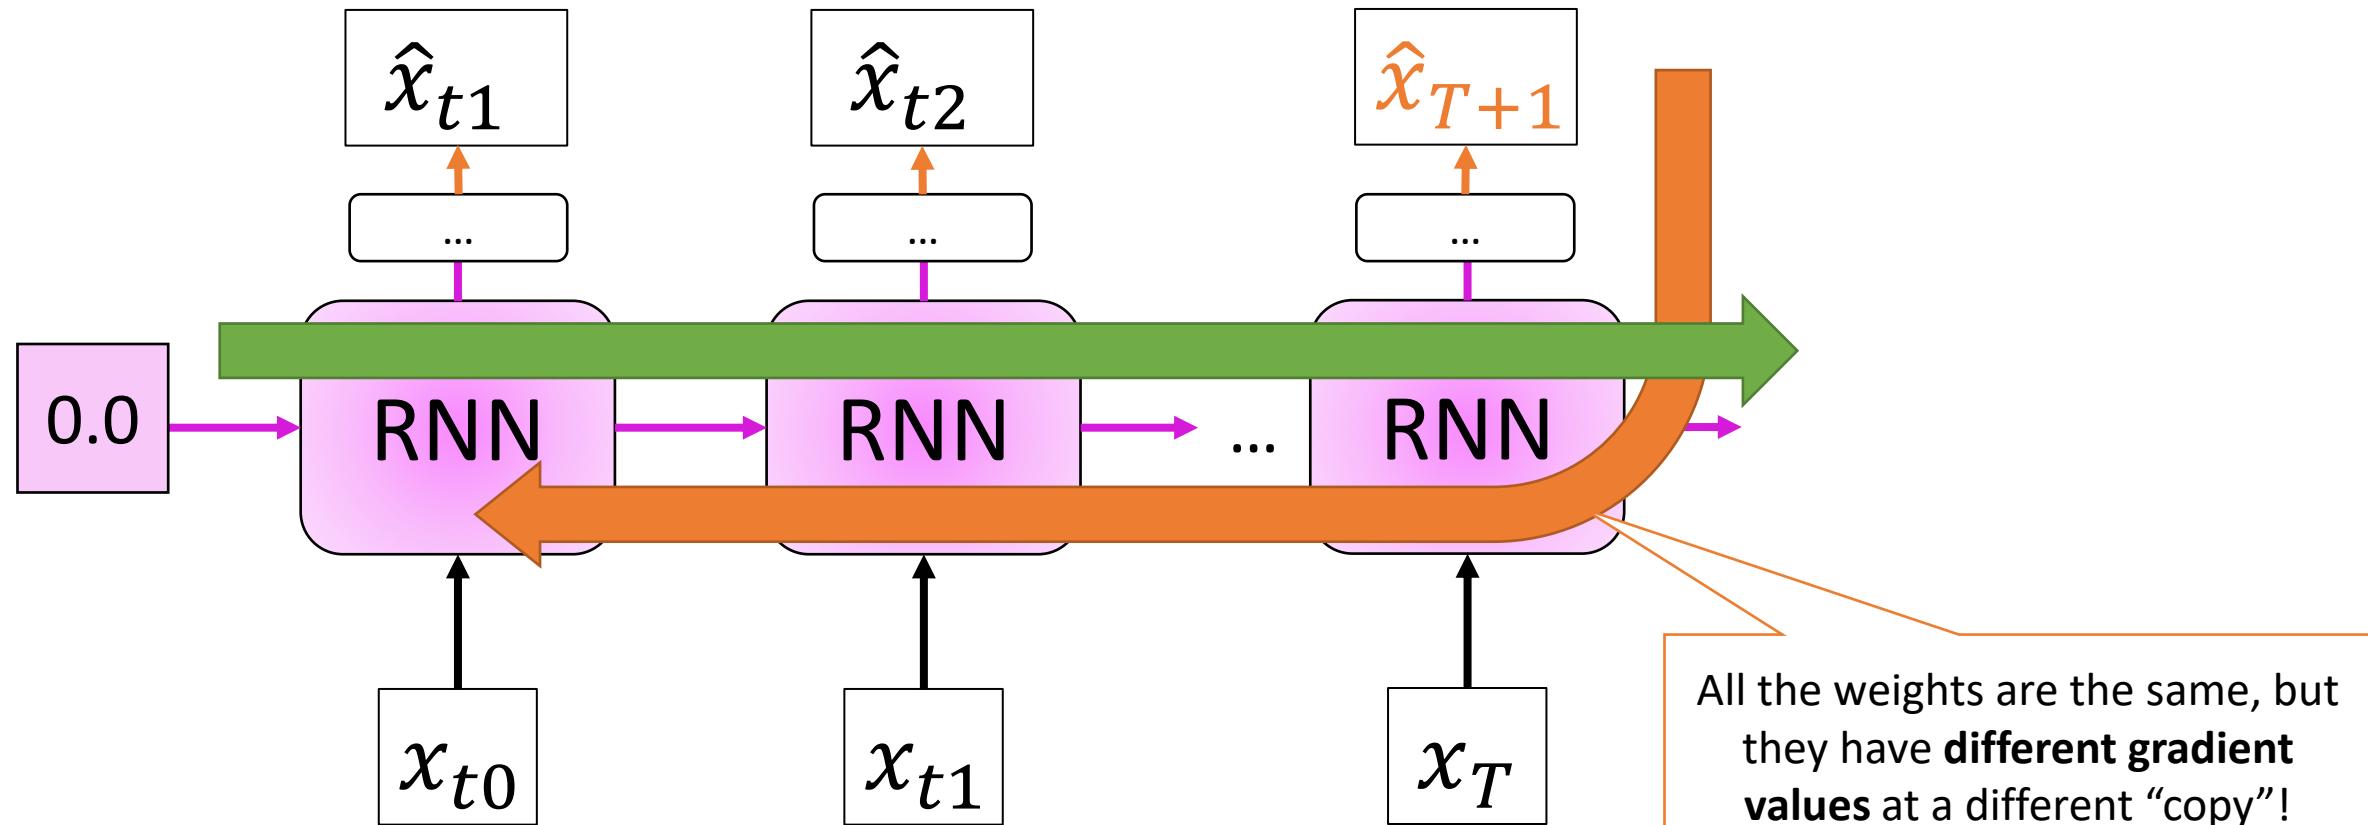
# > Recurrent Neural Networks

- How is a RNN trained? Unrolling!



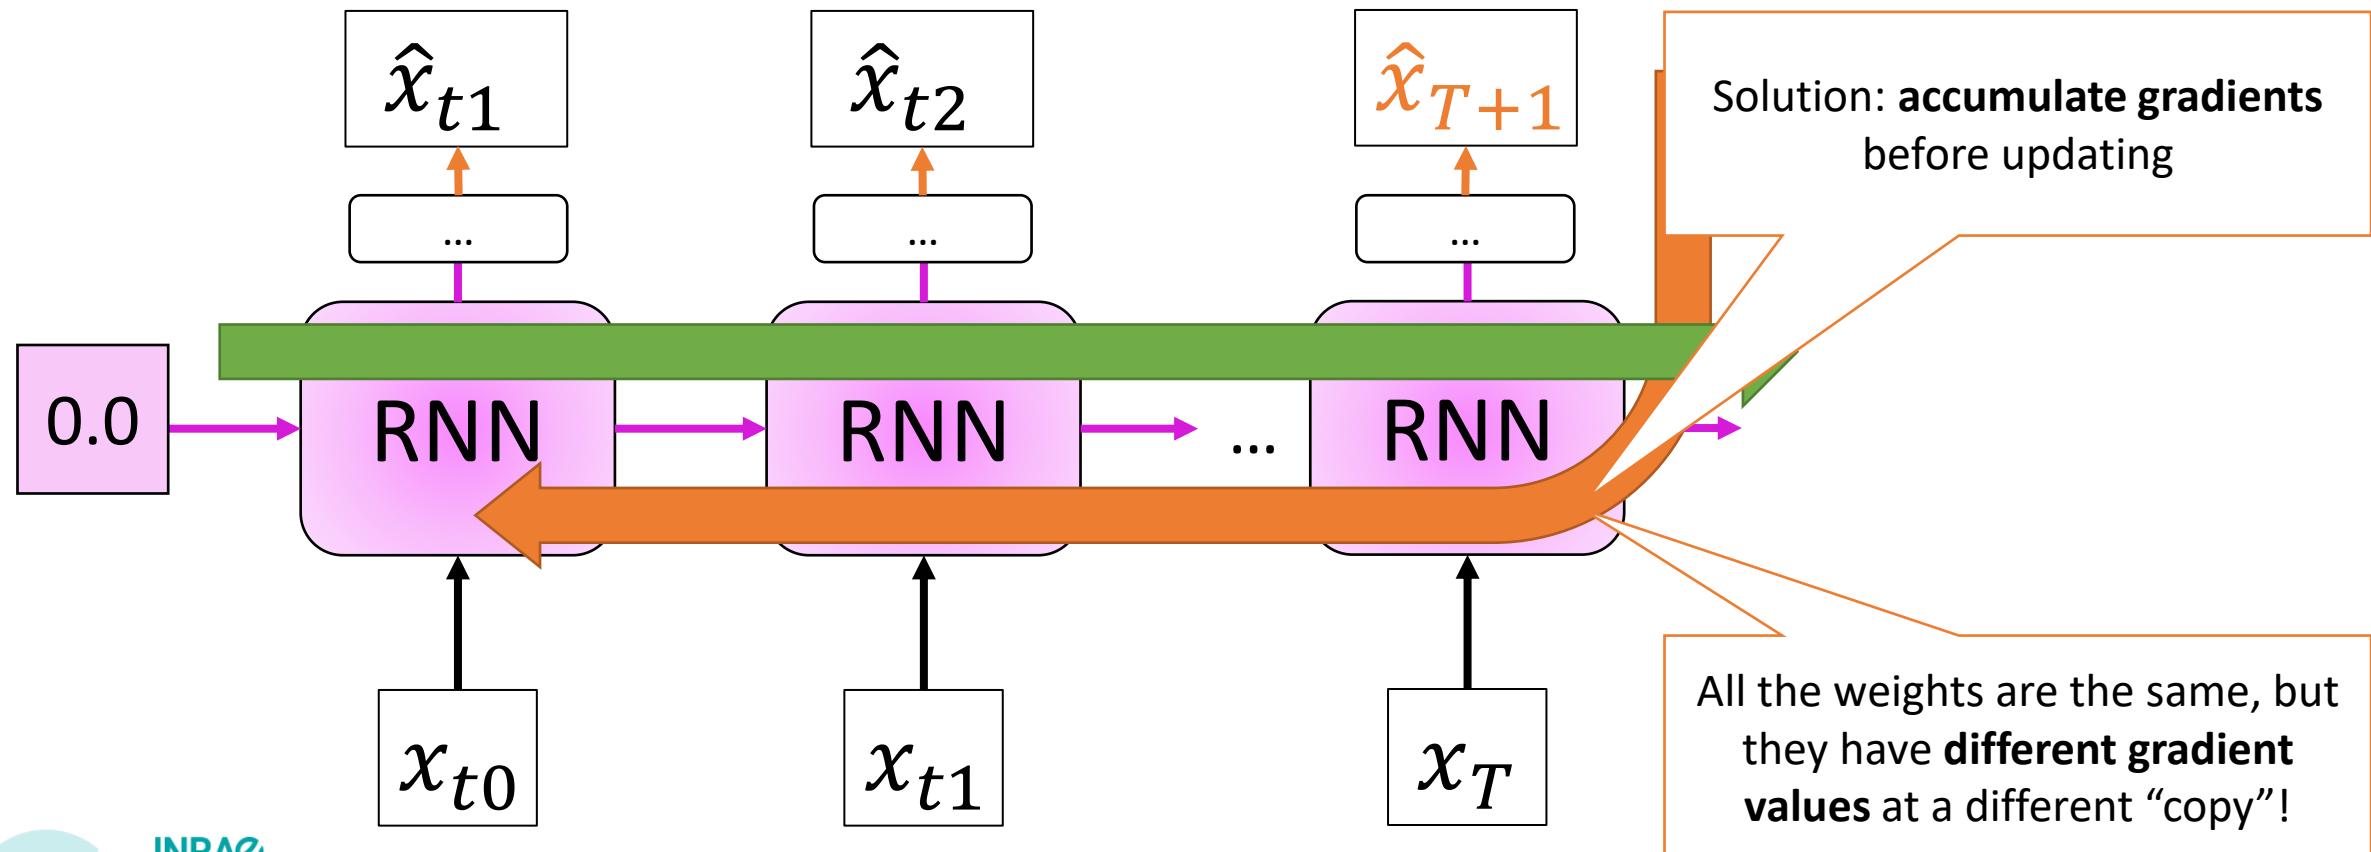
# > Recurrent Neural Networks

- How is a RNN trained? Unrolling!

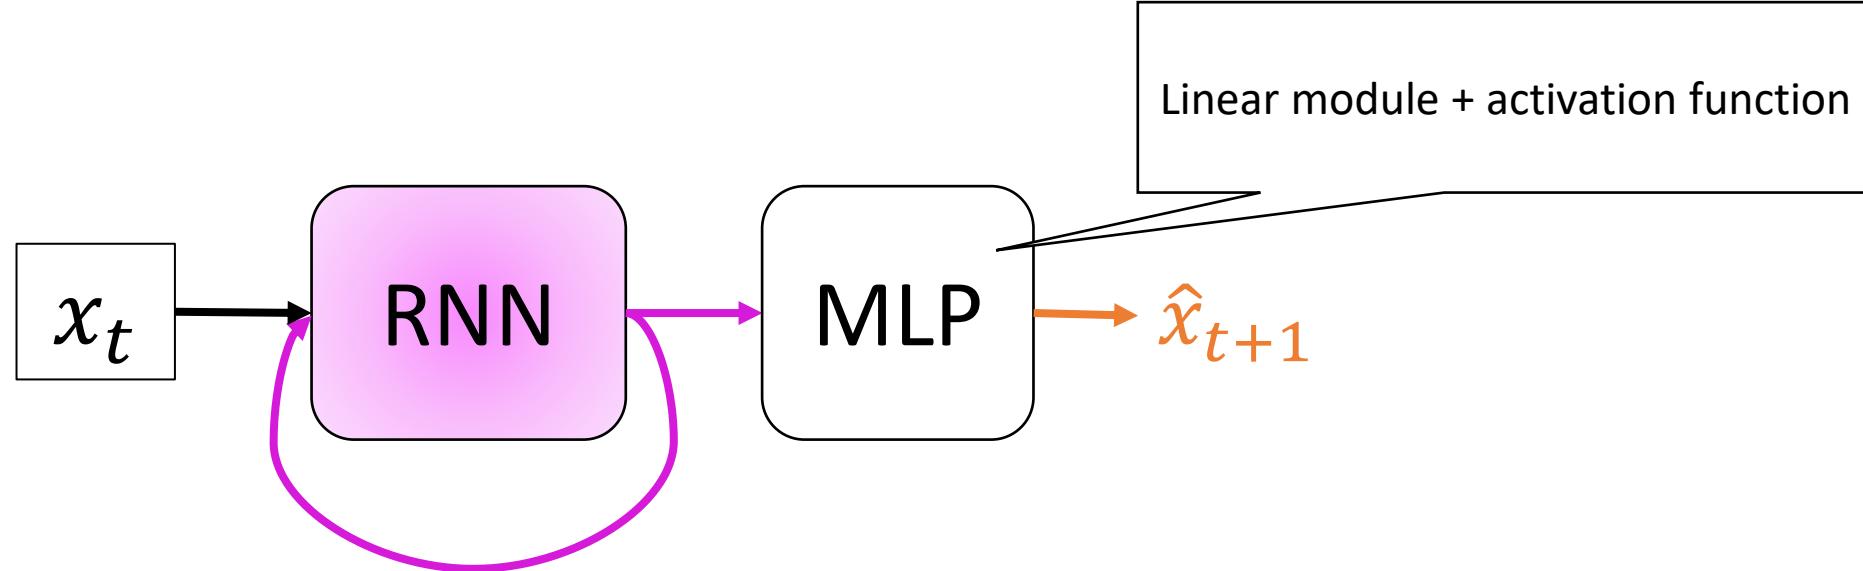


# > Recurrent Neural Networks

- How is a RNN trained? Unrolling!



# > Recurrent Neural Networks



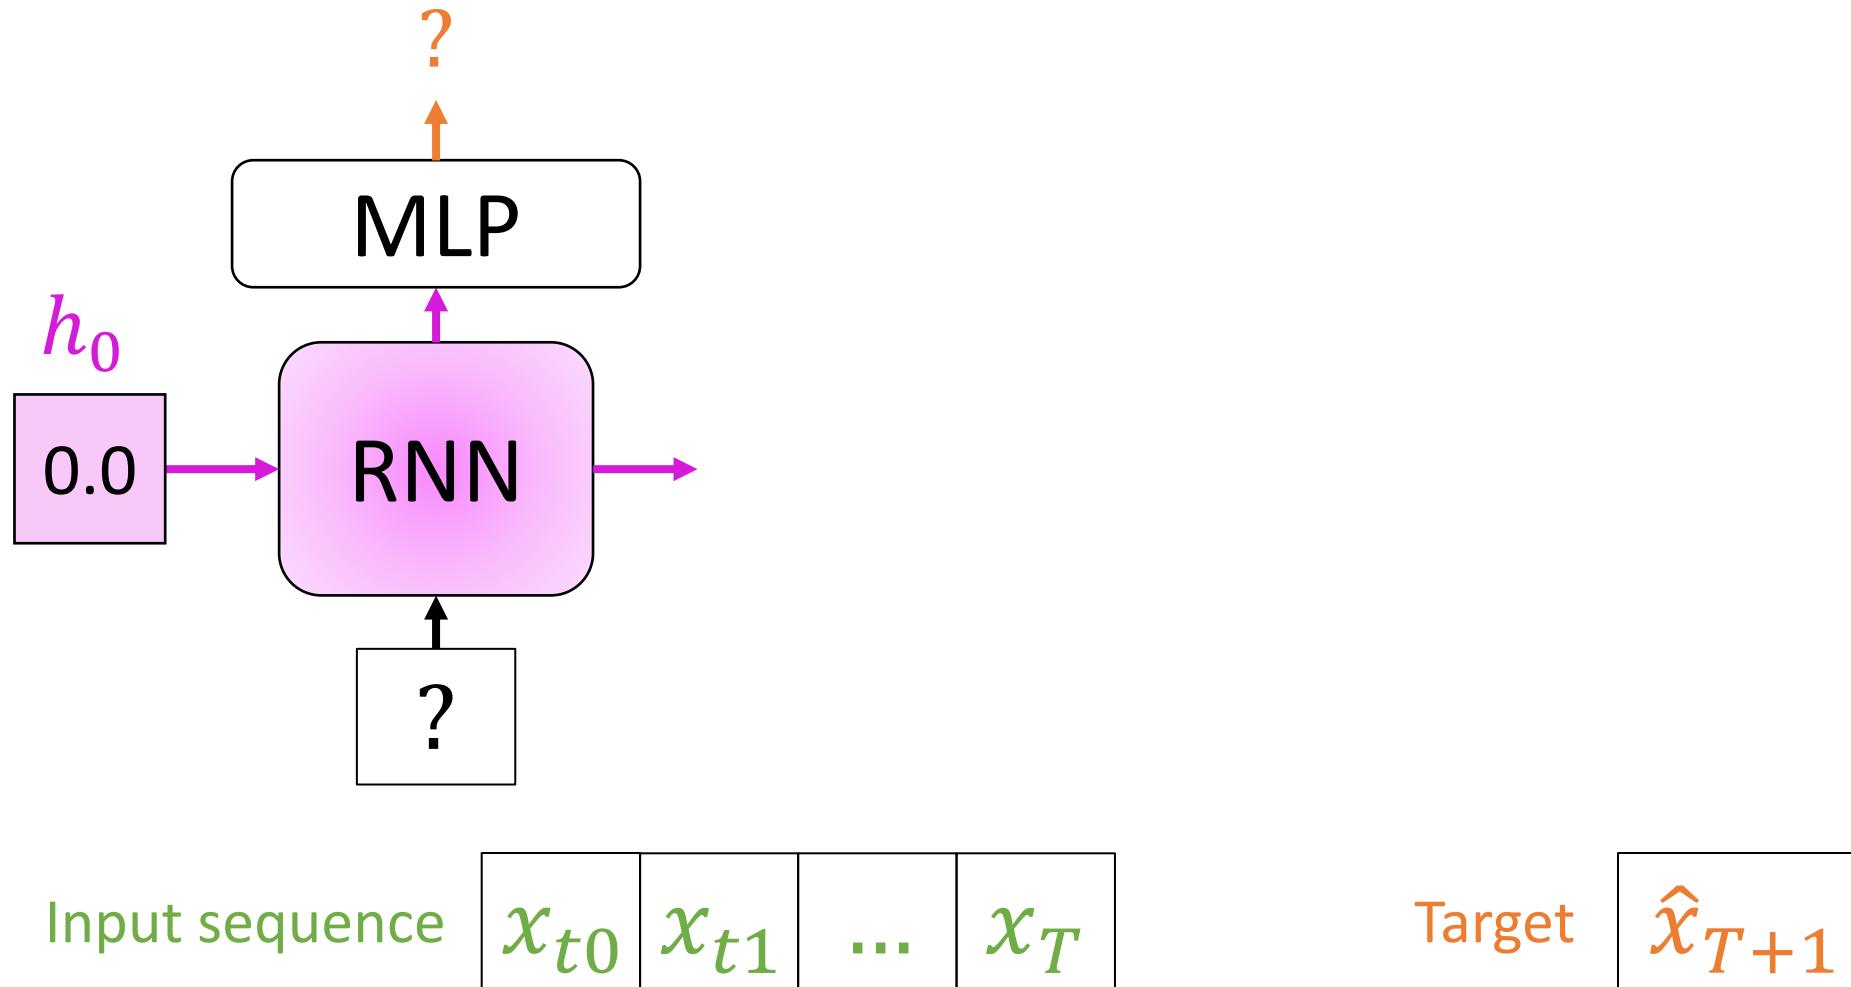
Input sequence

$$\begin{array}{|c|c|c|c|} \hline x_{t0} & x_{t1} & \dots & x_T \\ \hline \end{array}$$

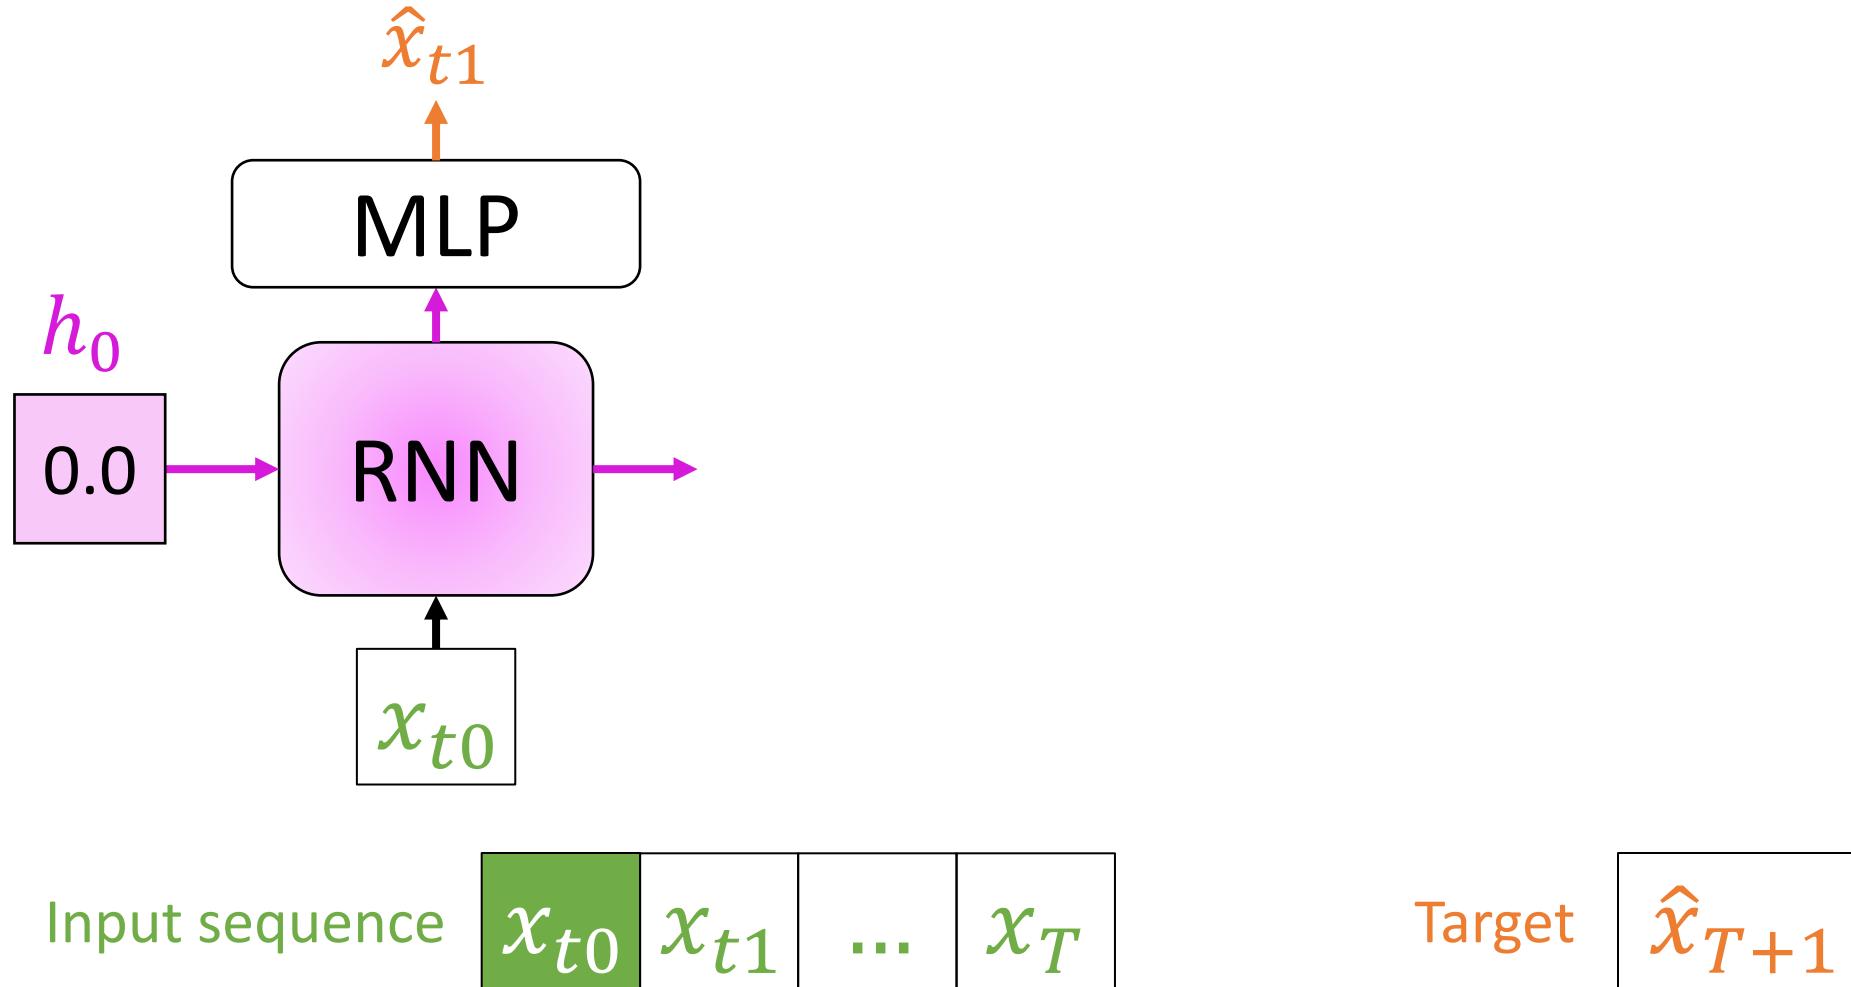
Target

$$\hat{x}_{T+1}$$

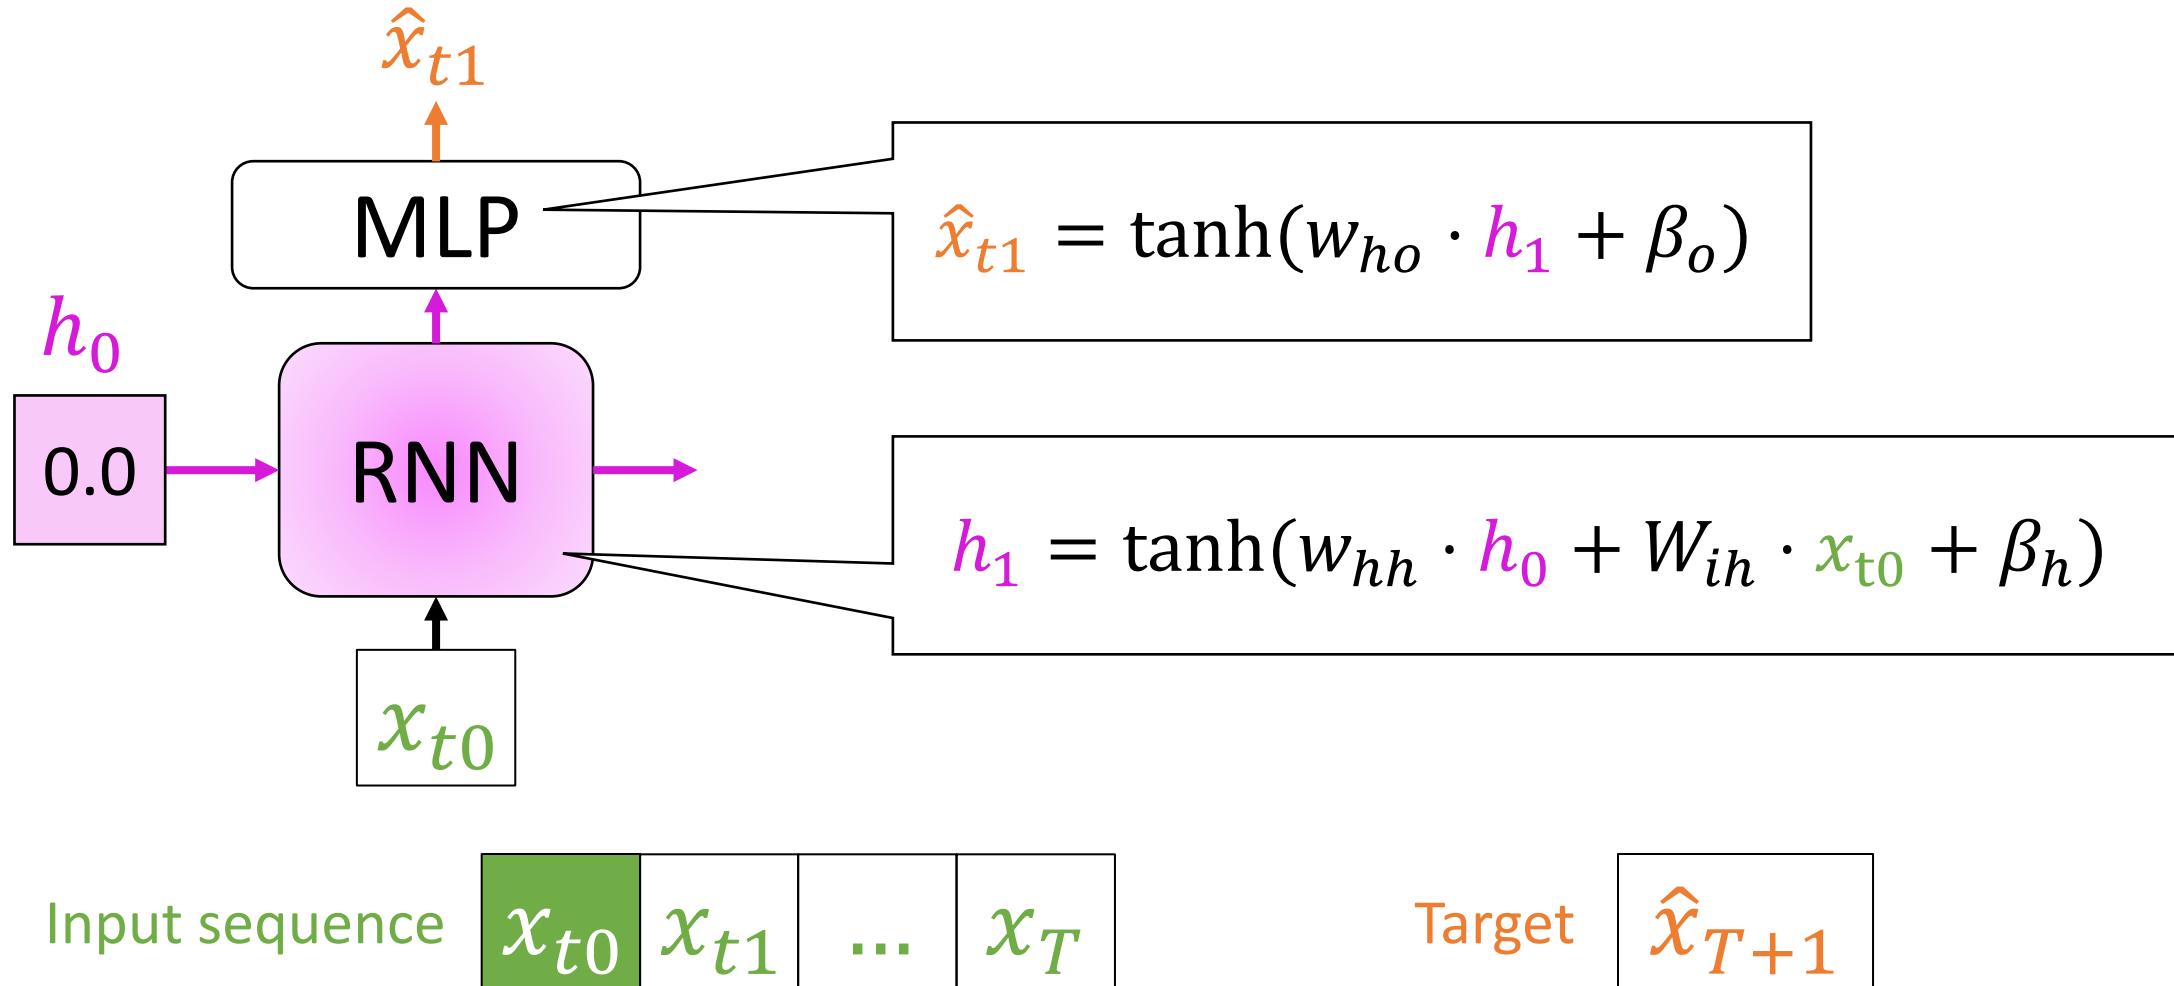
# > Recurrent Neural Networks



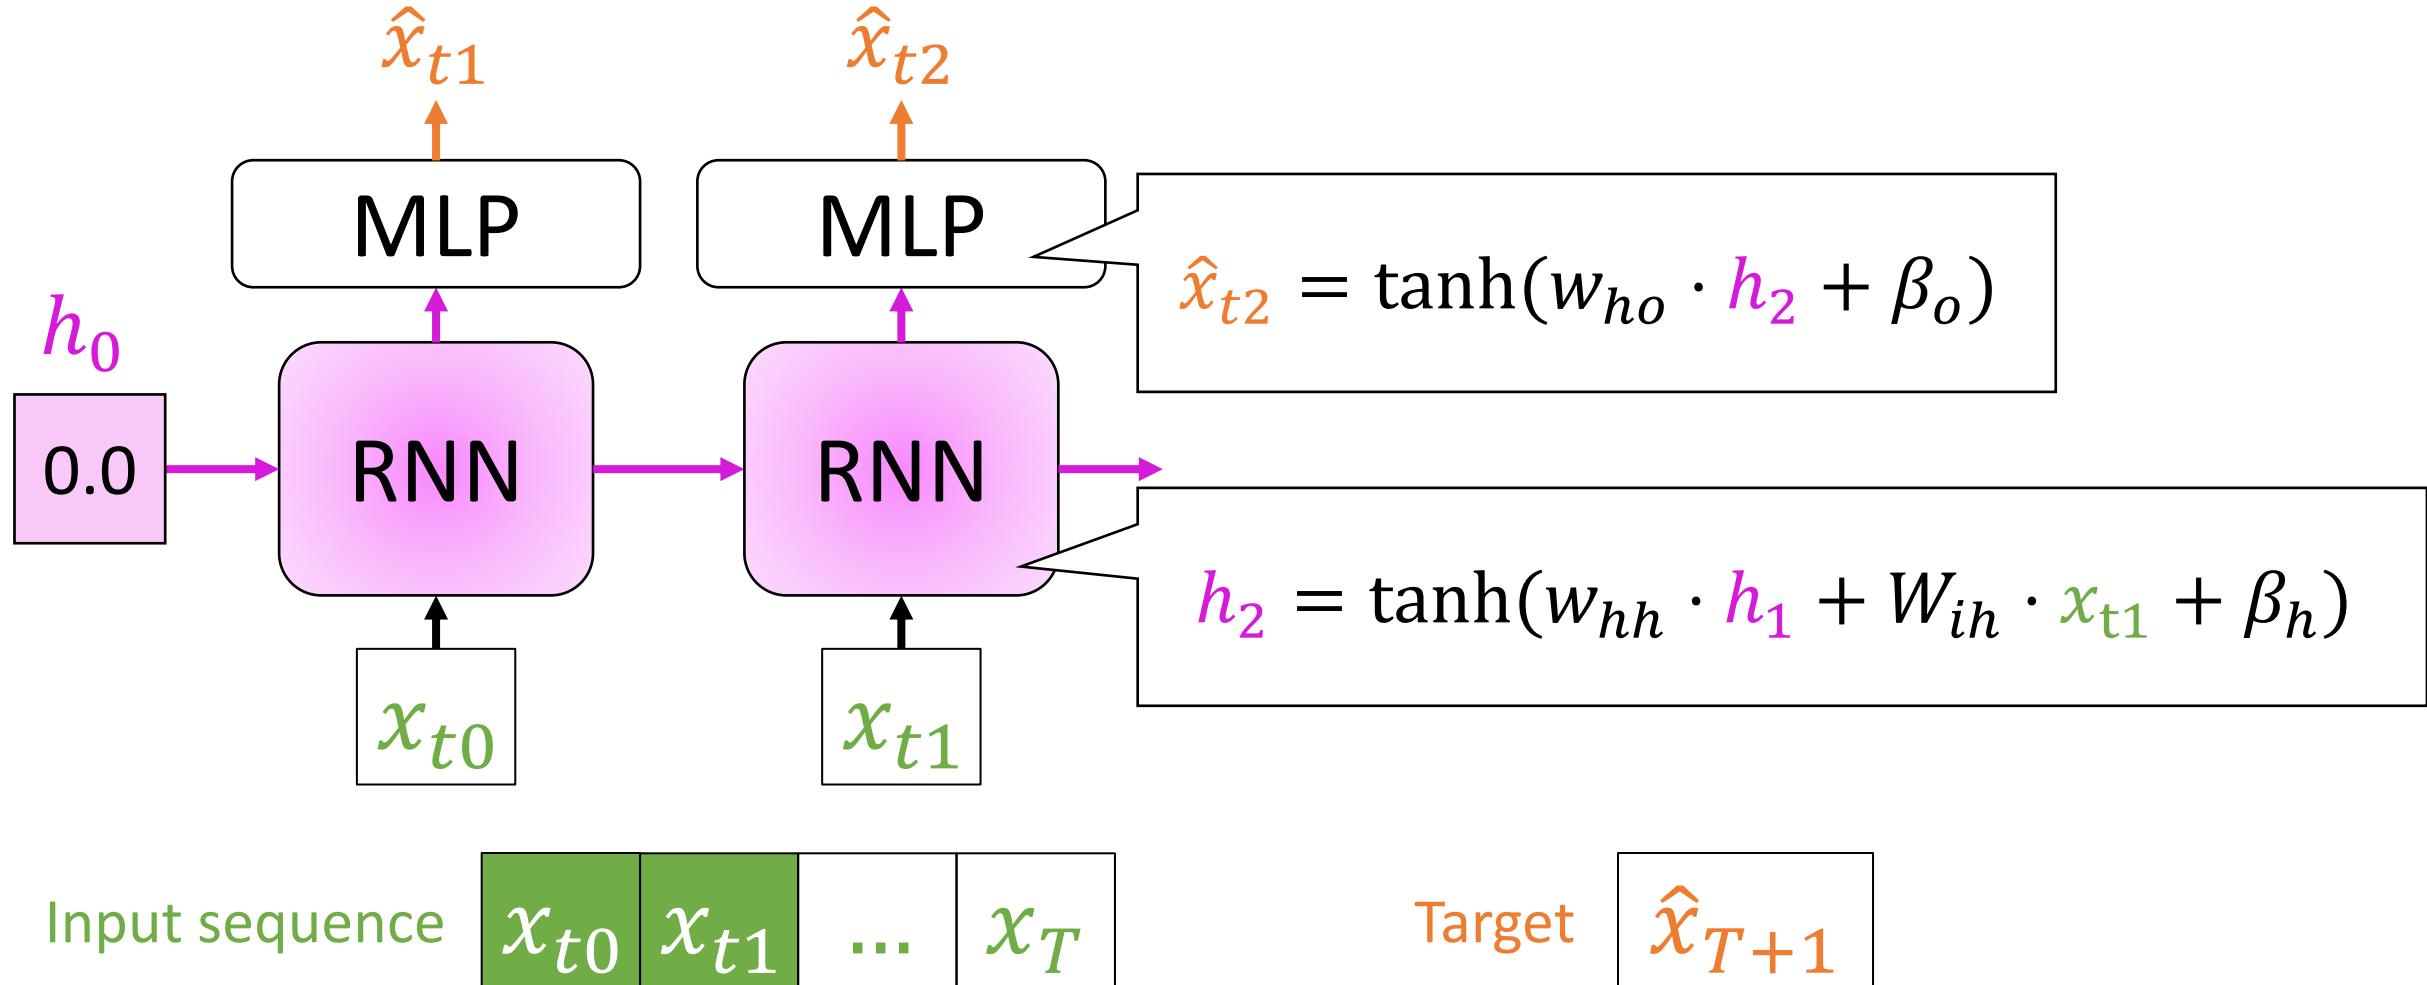
# > Recurrent Neural Networks



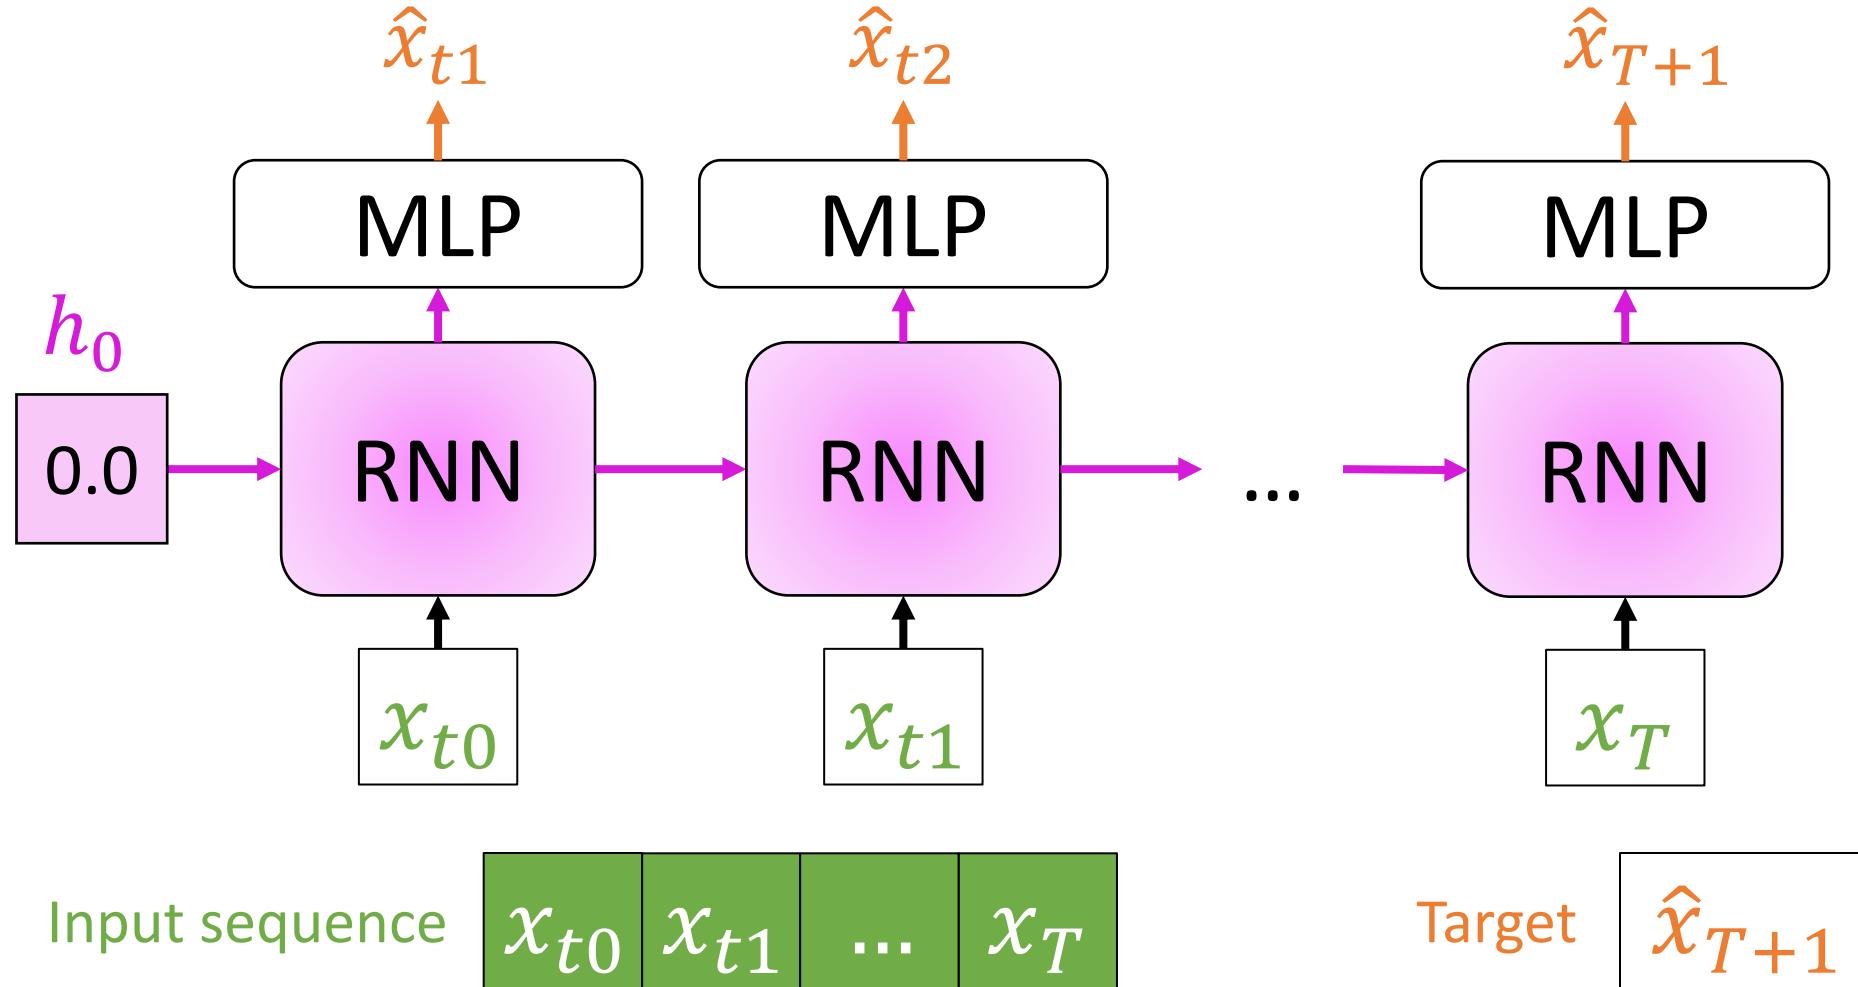
# > Recurrent Neural Networks



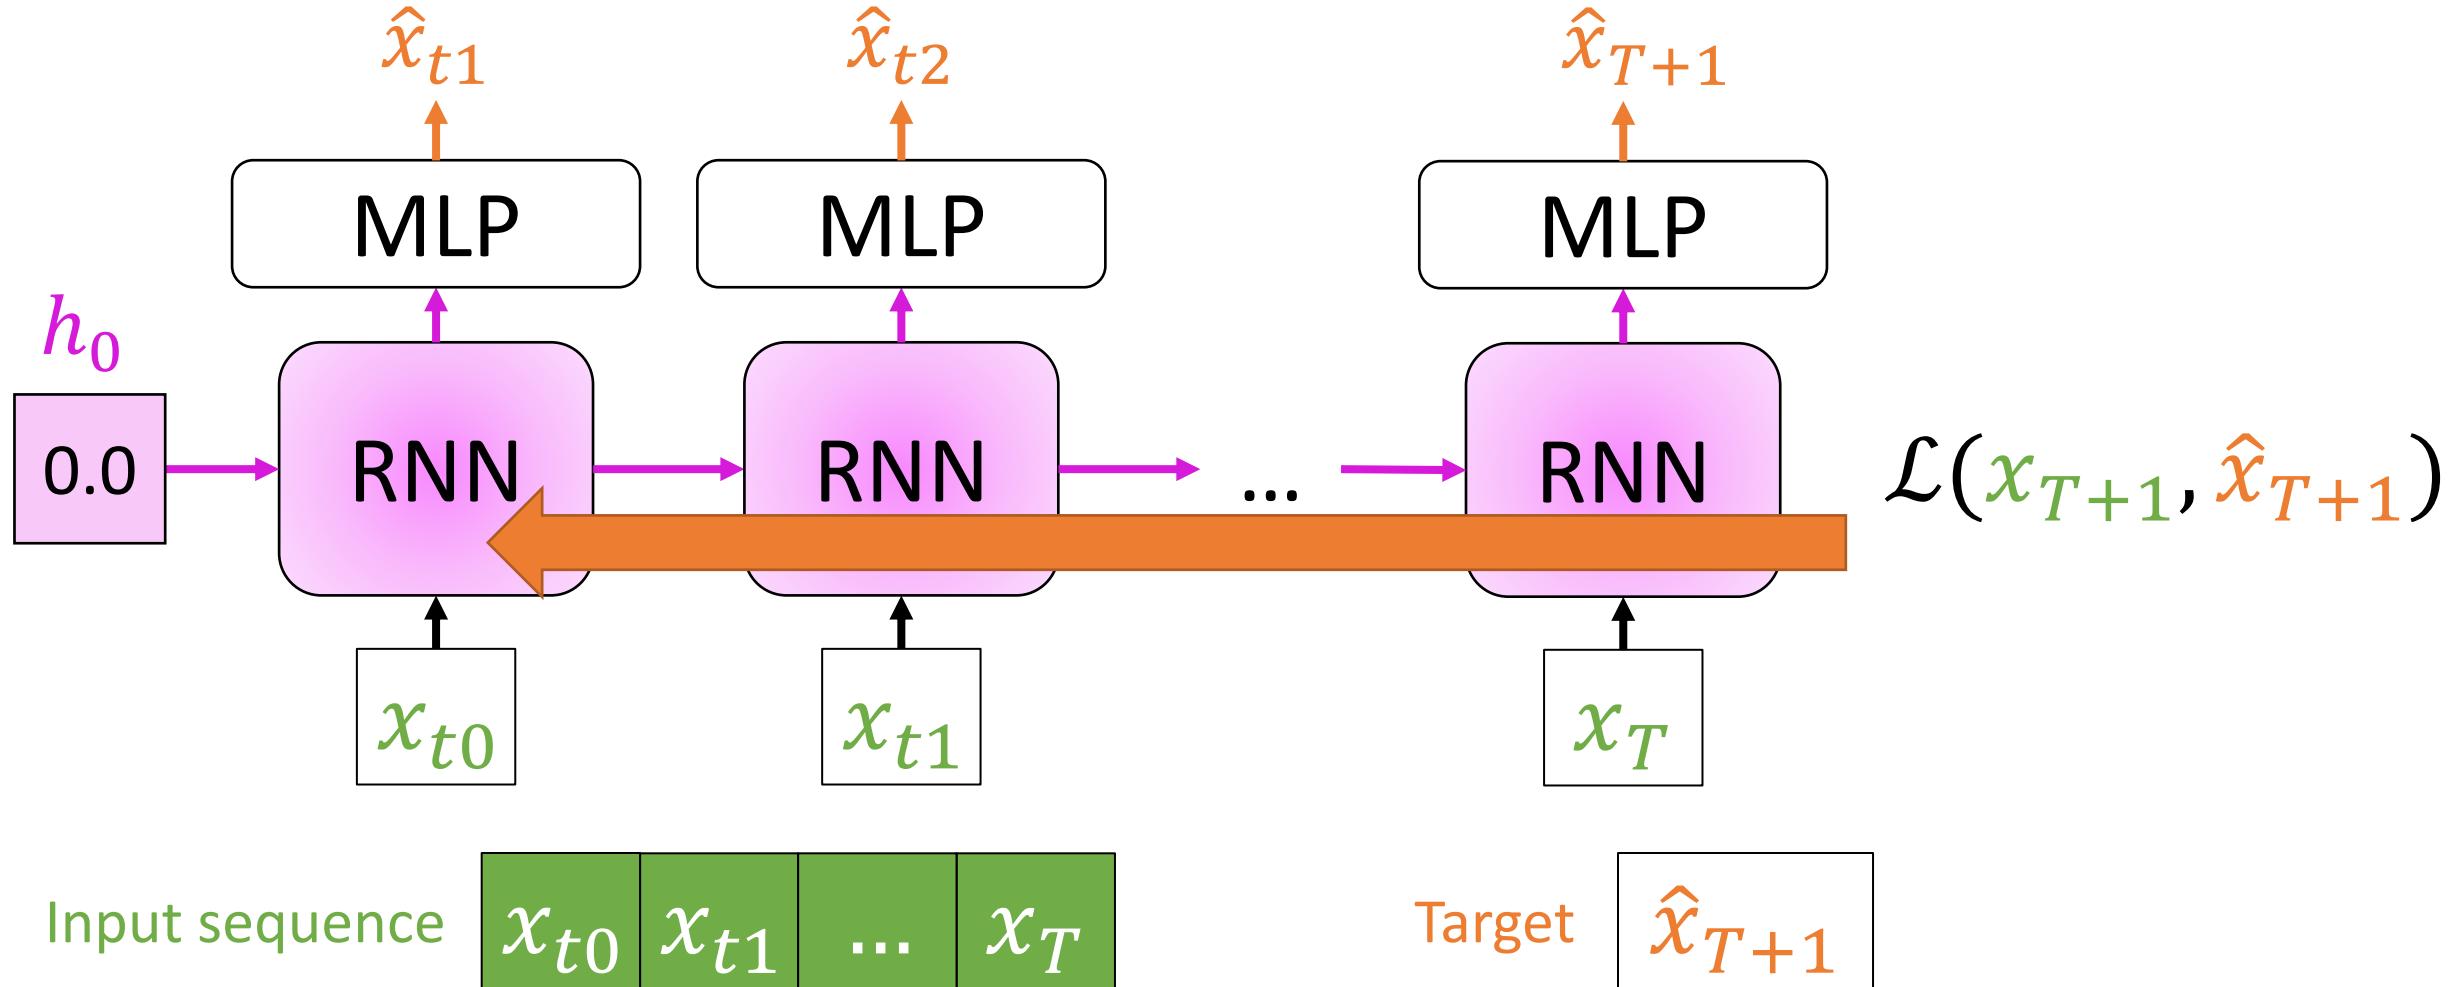
# > Recurrent Neural Networks



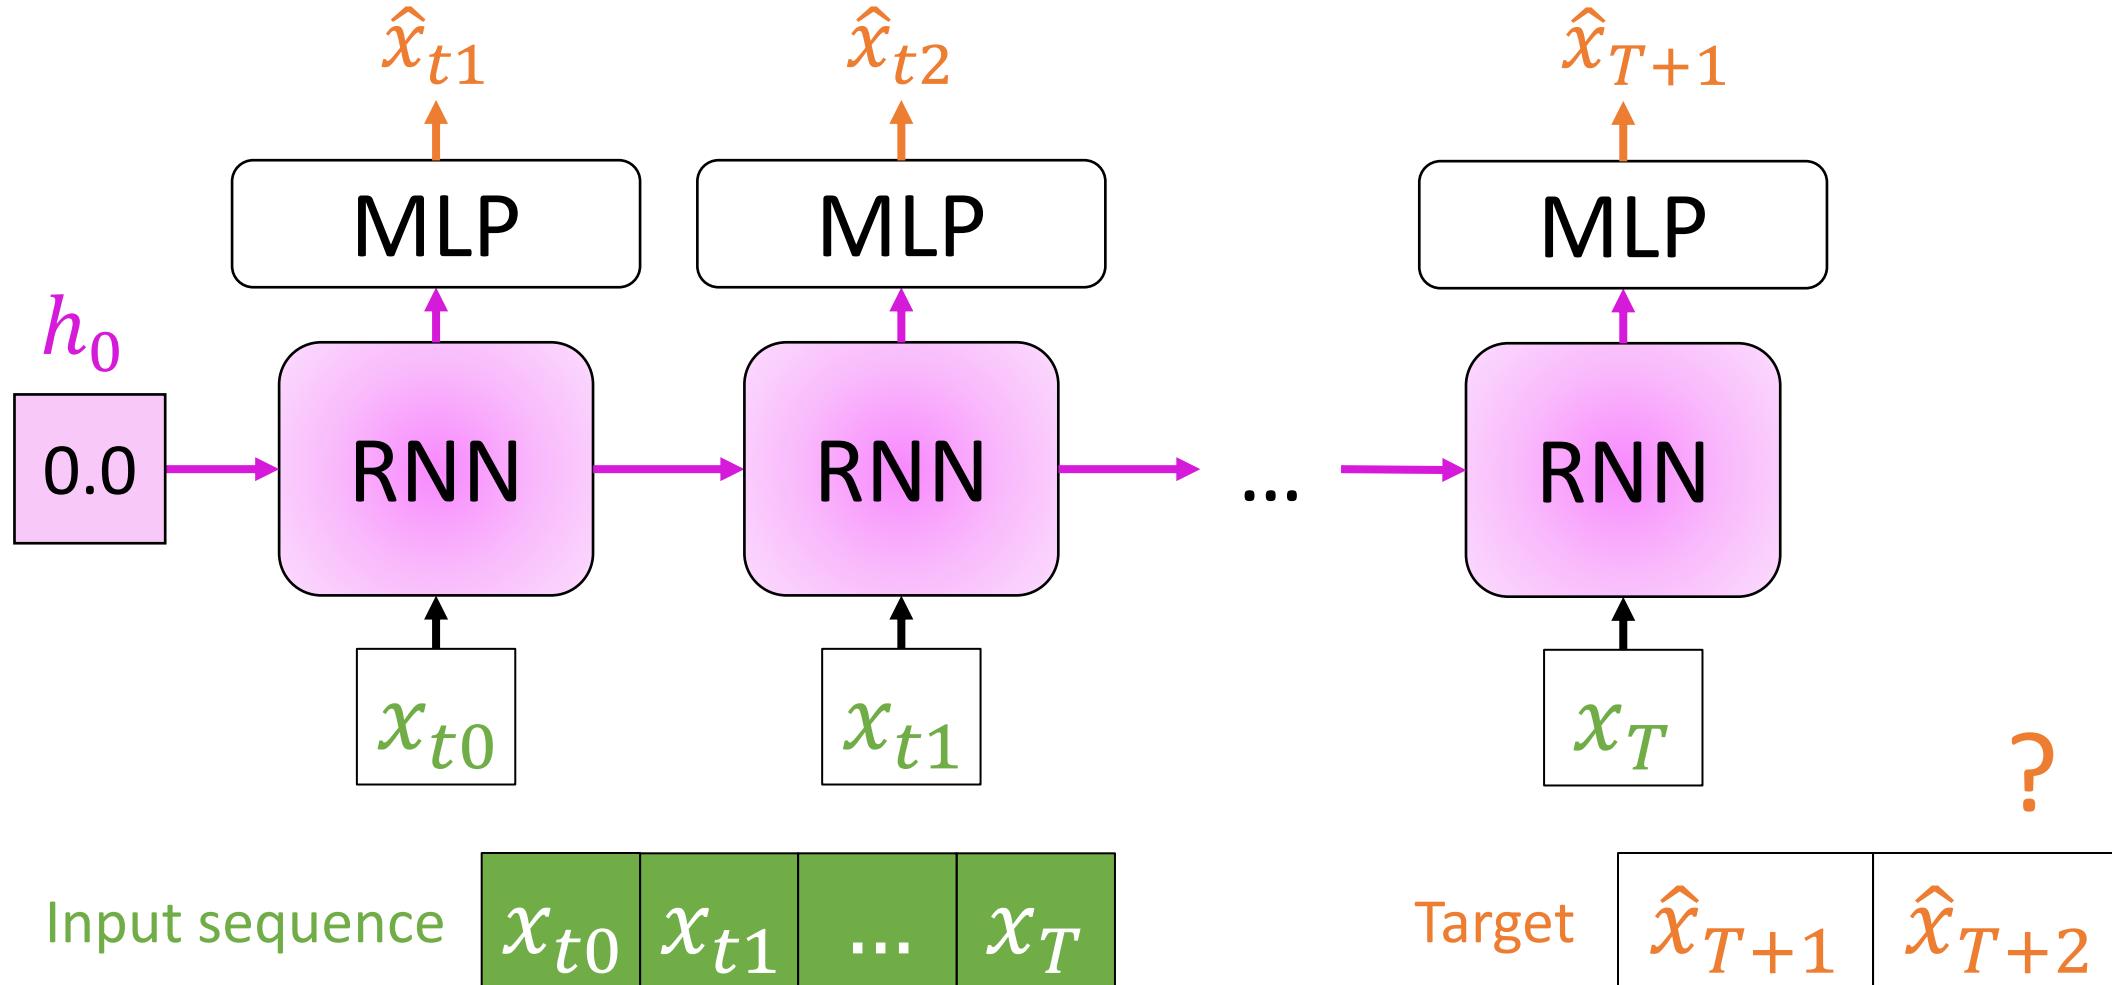
# > Recurrent Neural Networks



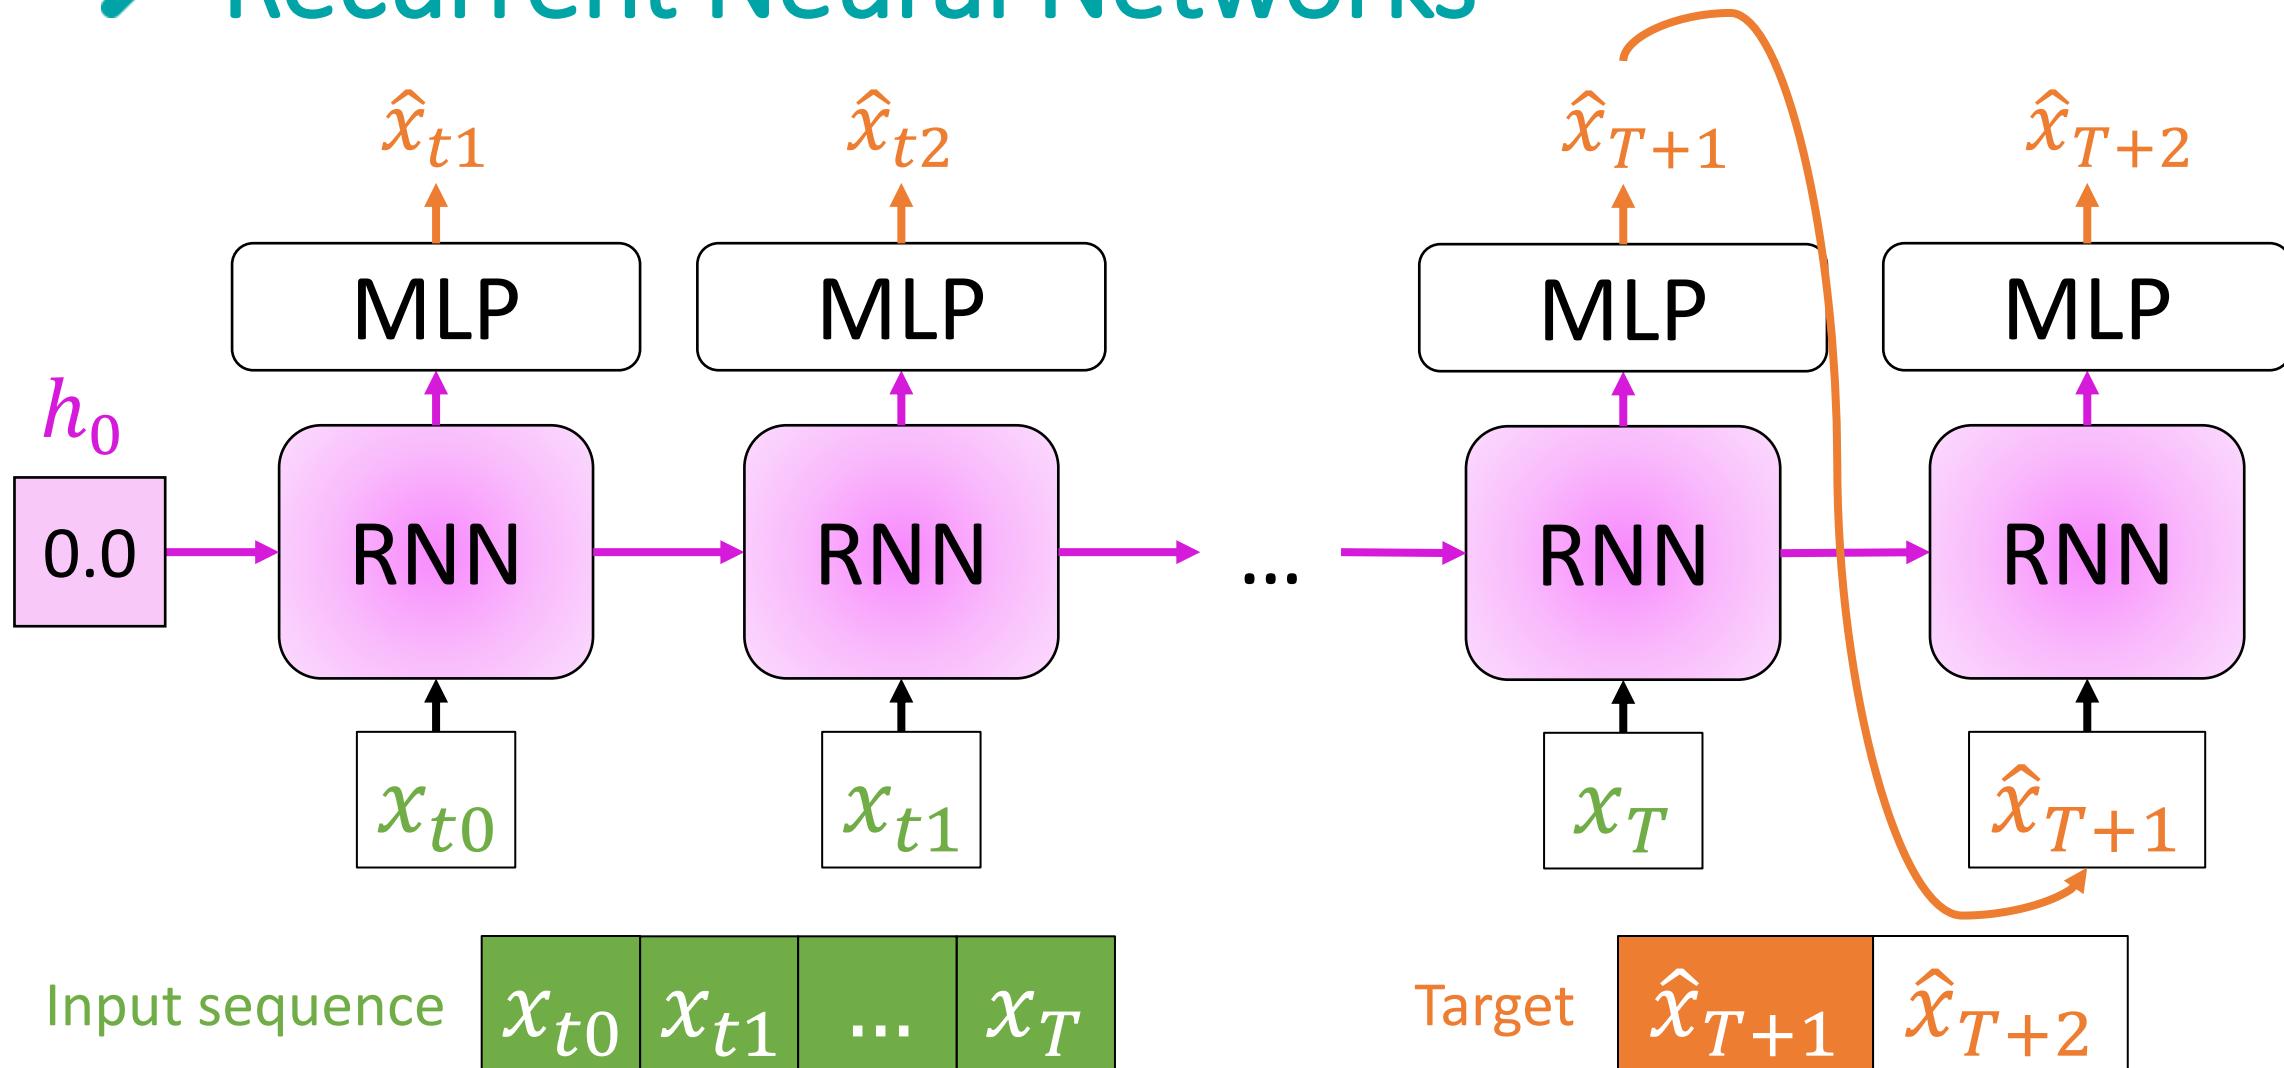
# > Recurrent Neural Networks



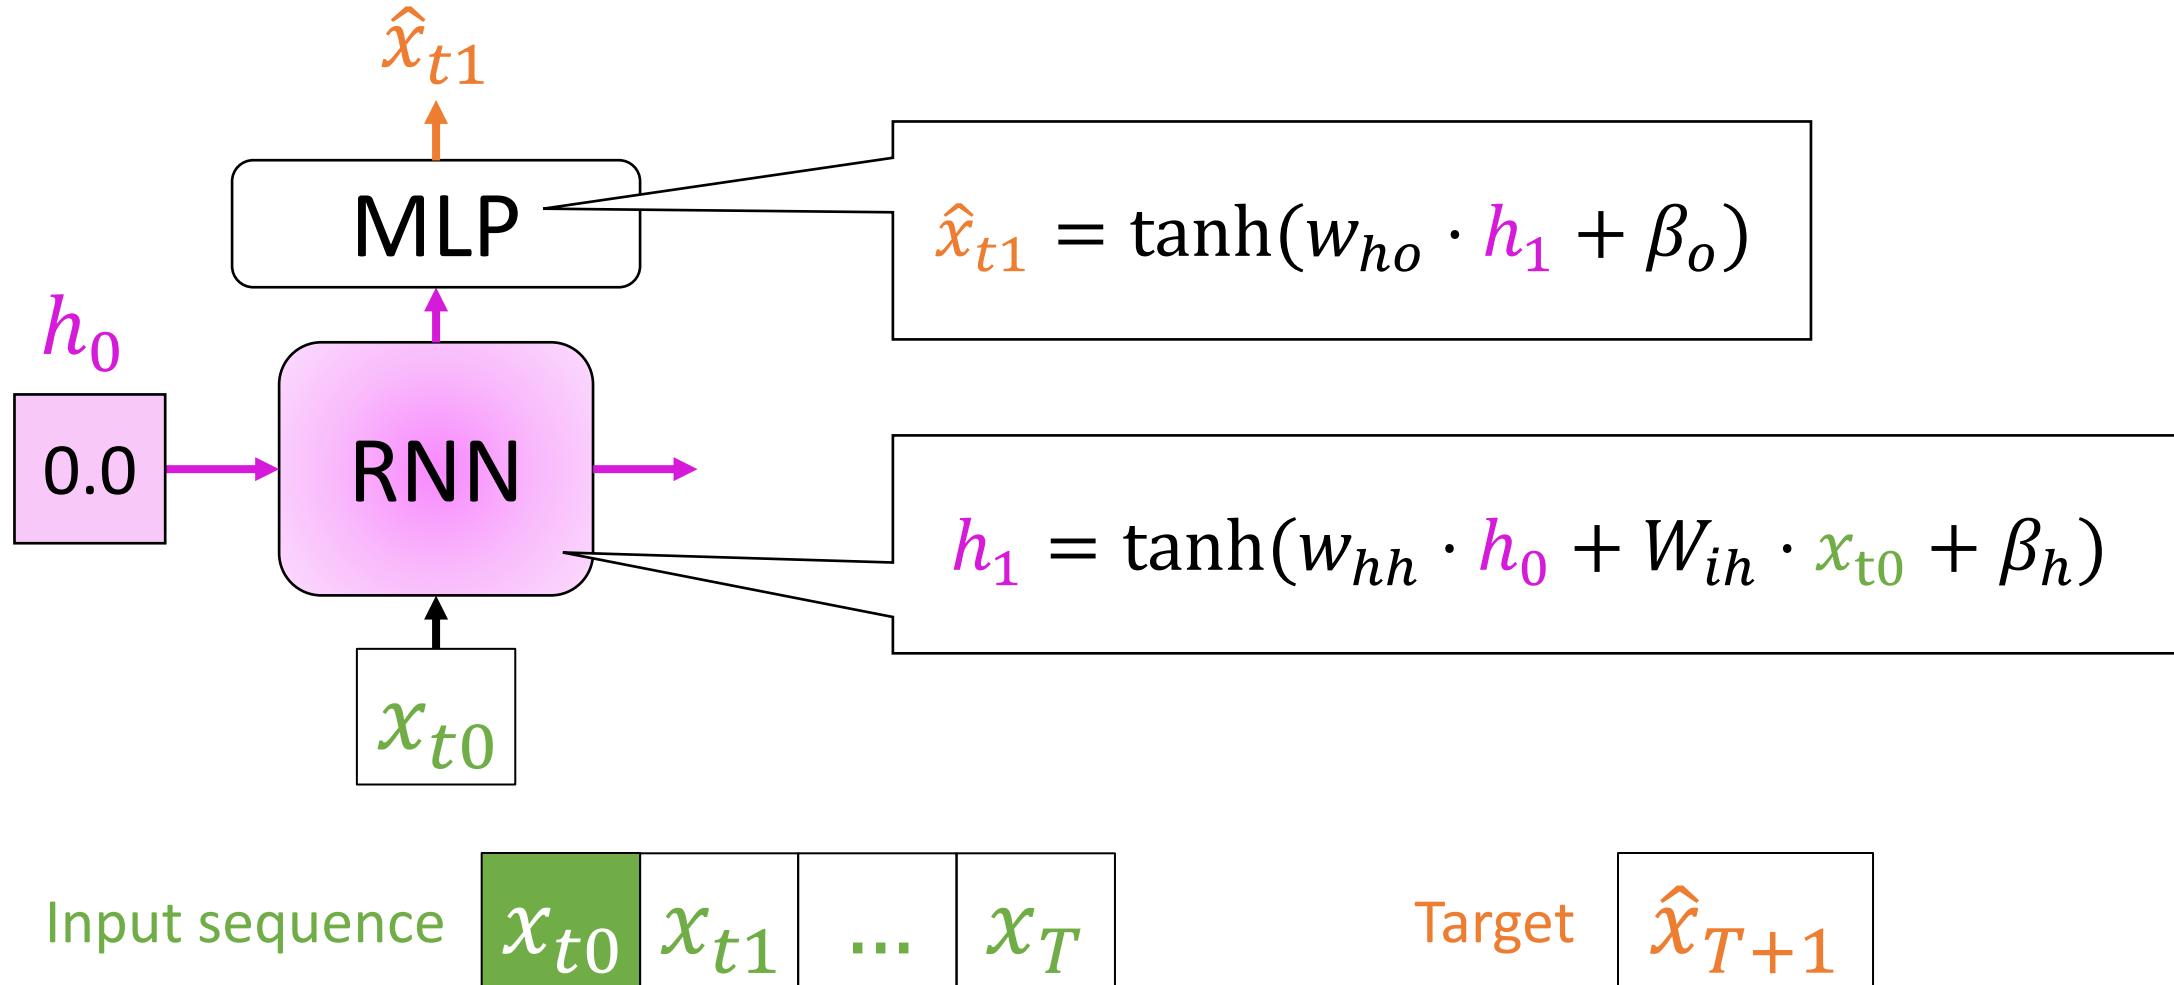
# > Recurrent Neural Networks



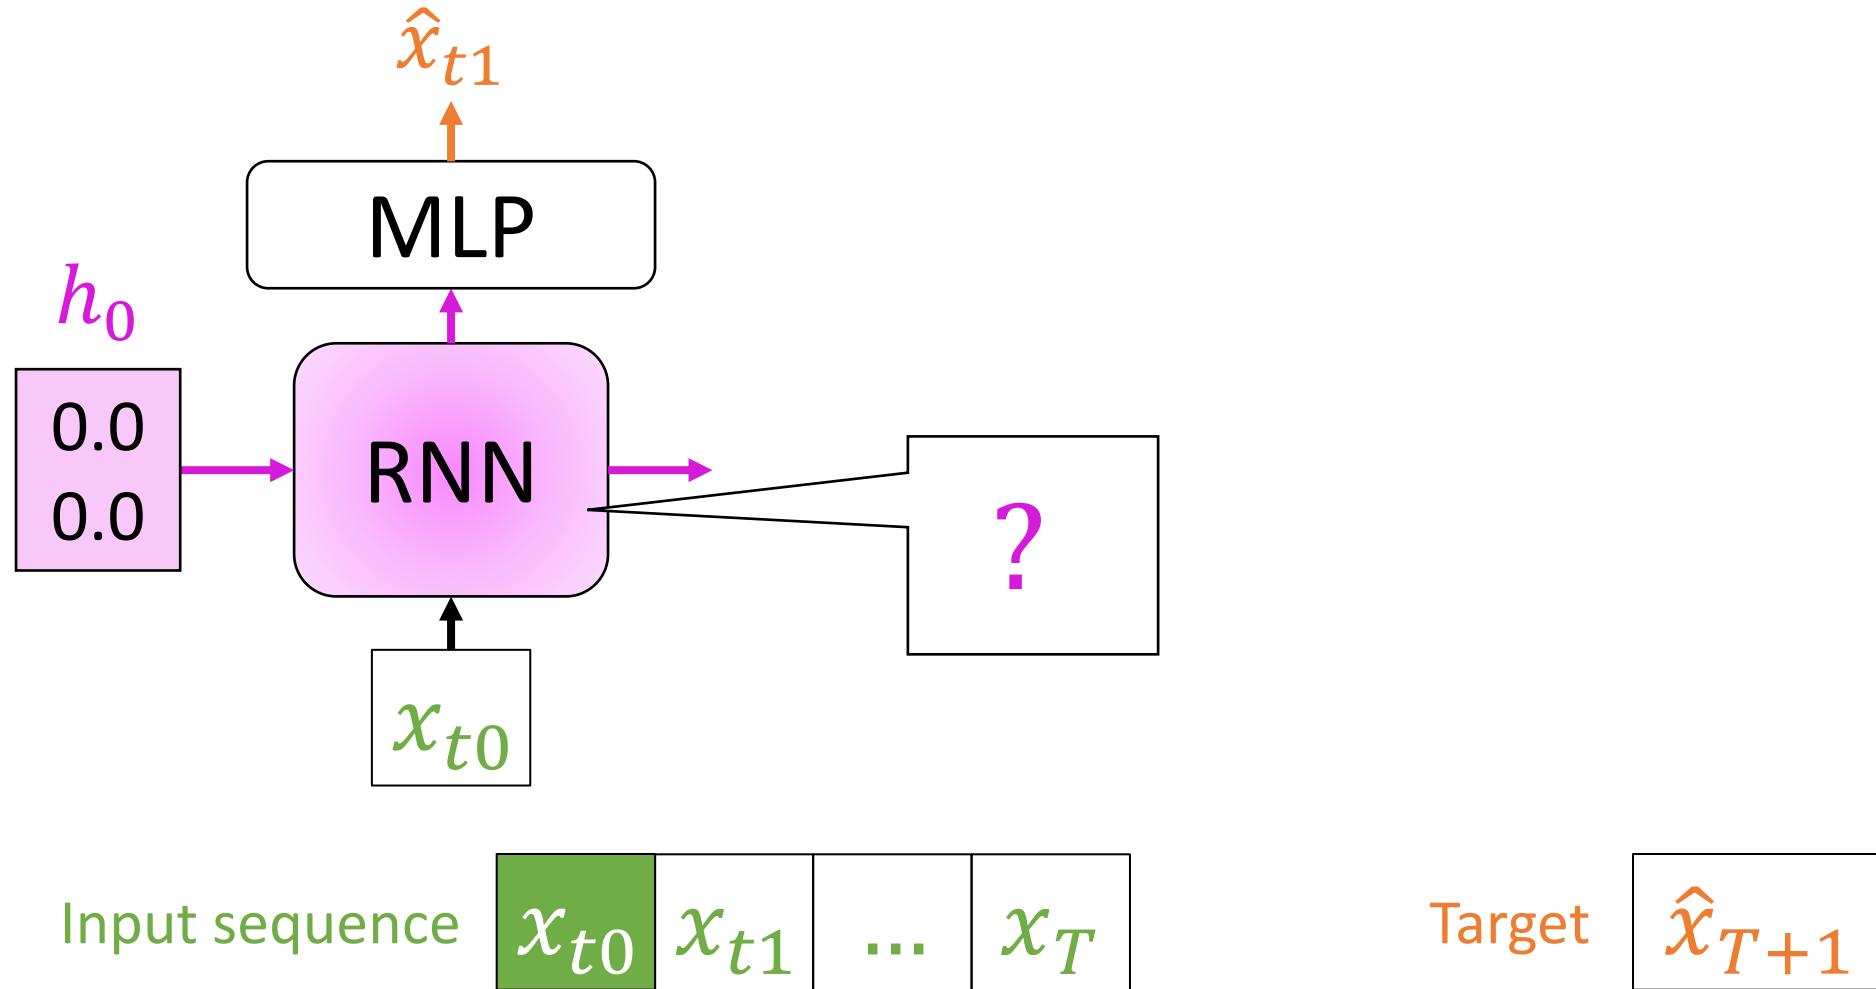
# > Recurrent Neural Networks



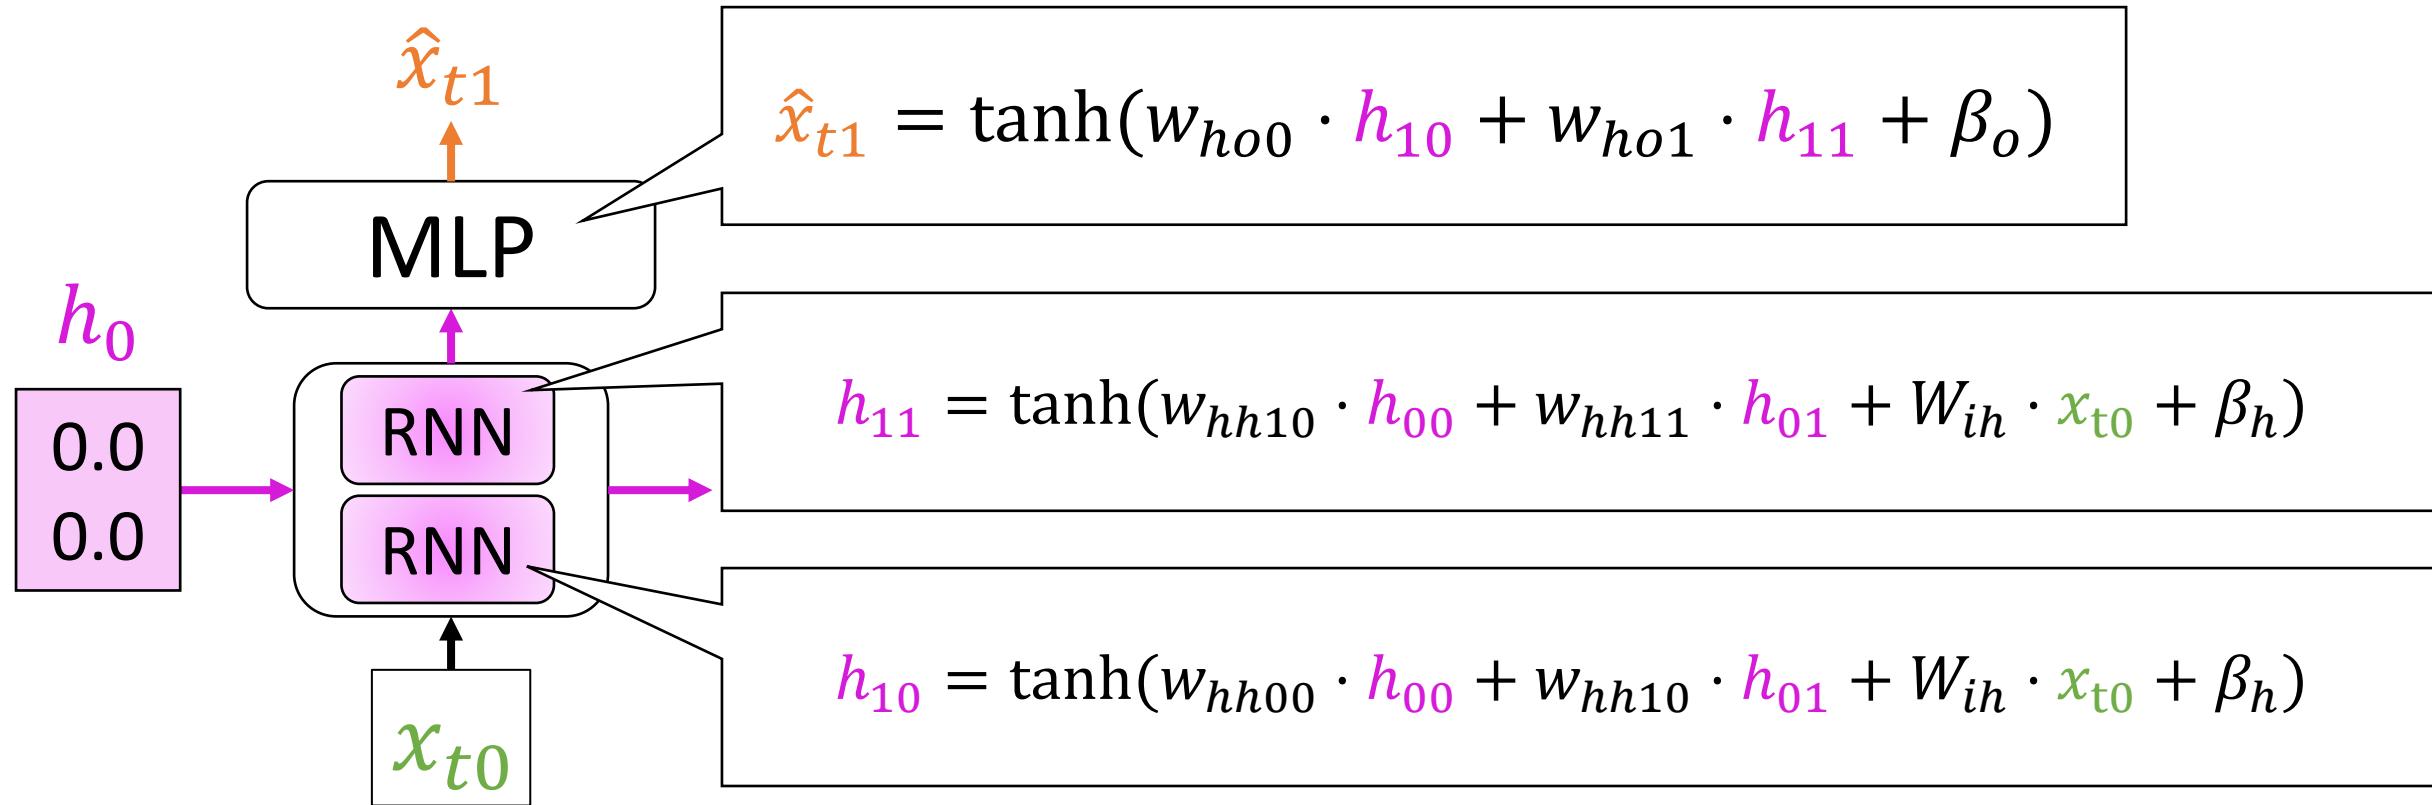
# > Recurrent Neural Networks



# > Recurrent Neural Networks



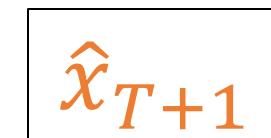
# > Recurrent Neural Networks



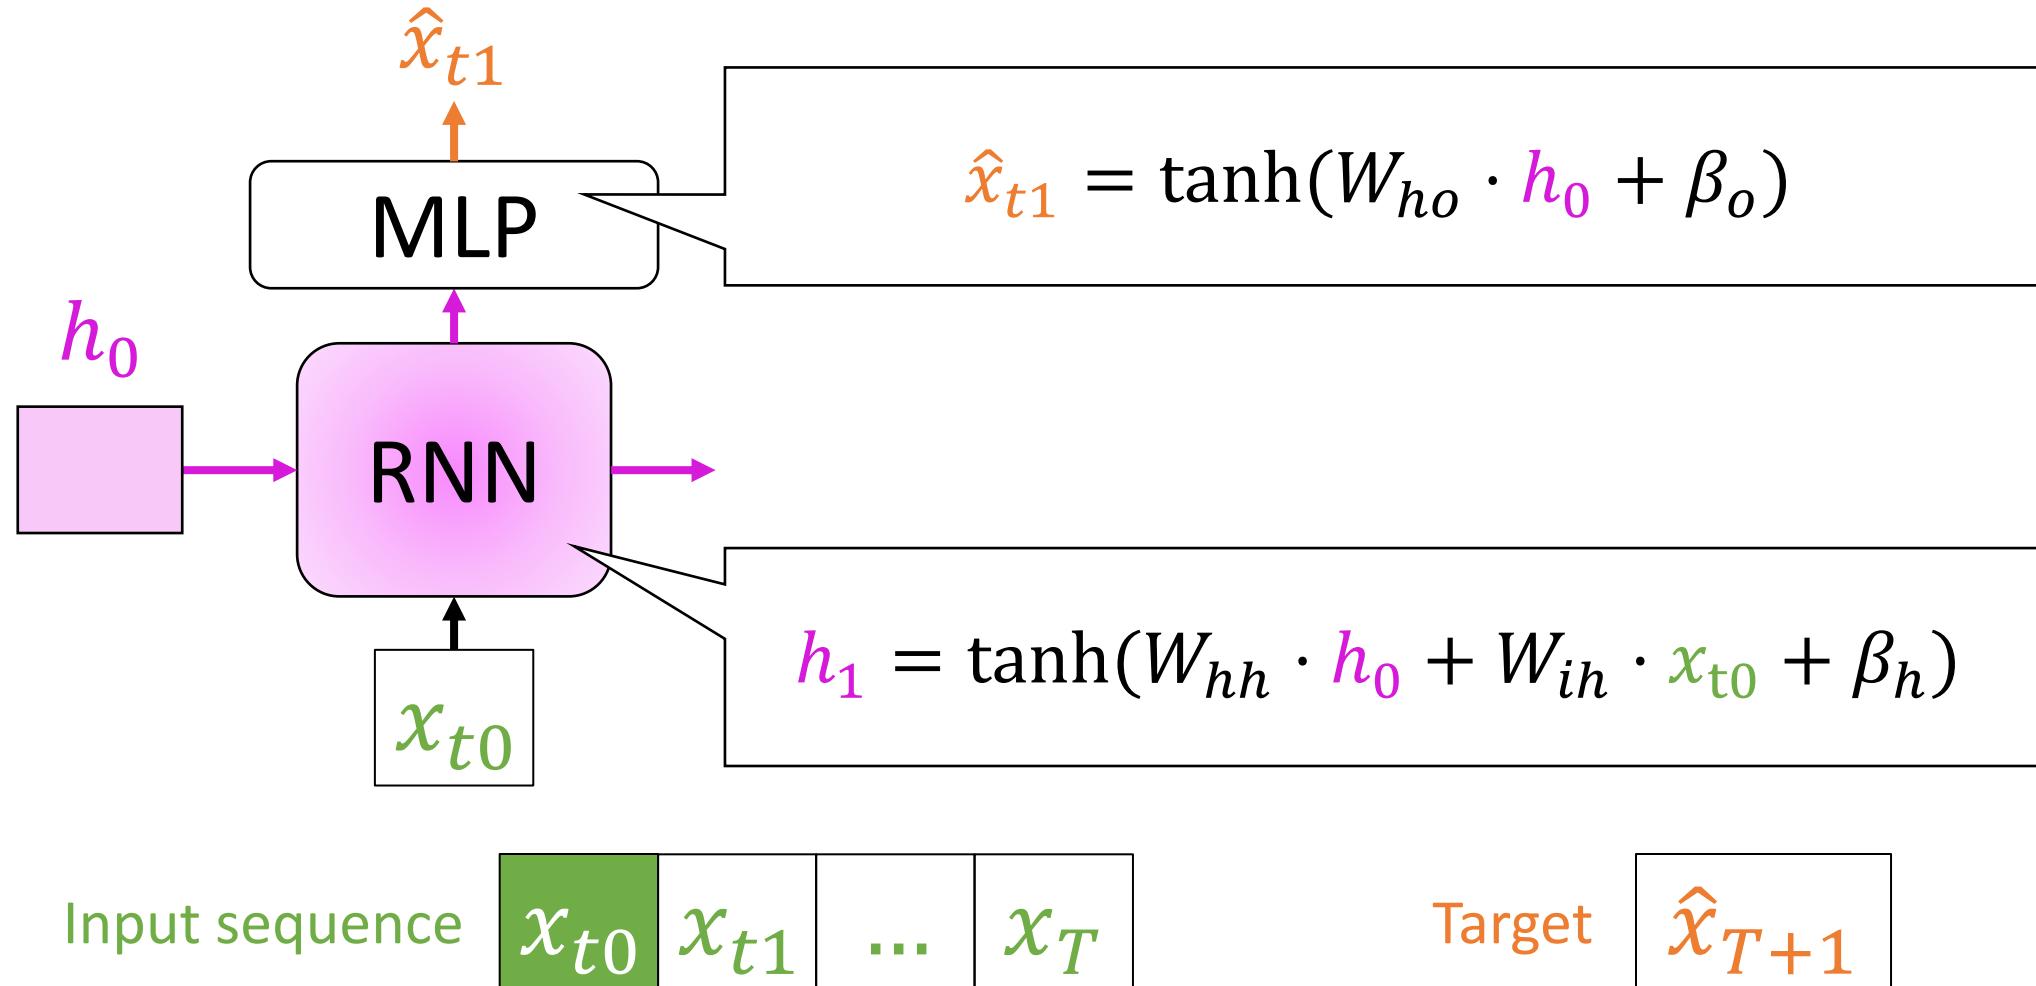
Input sequence



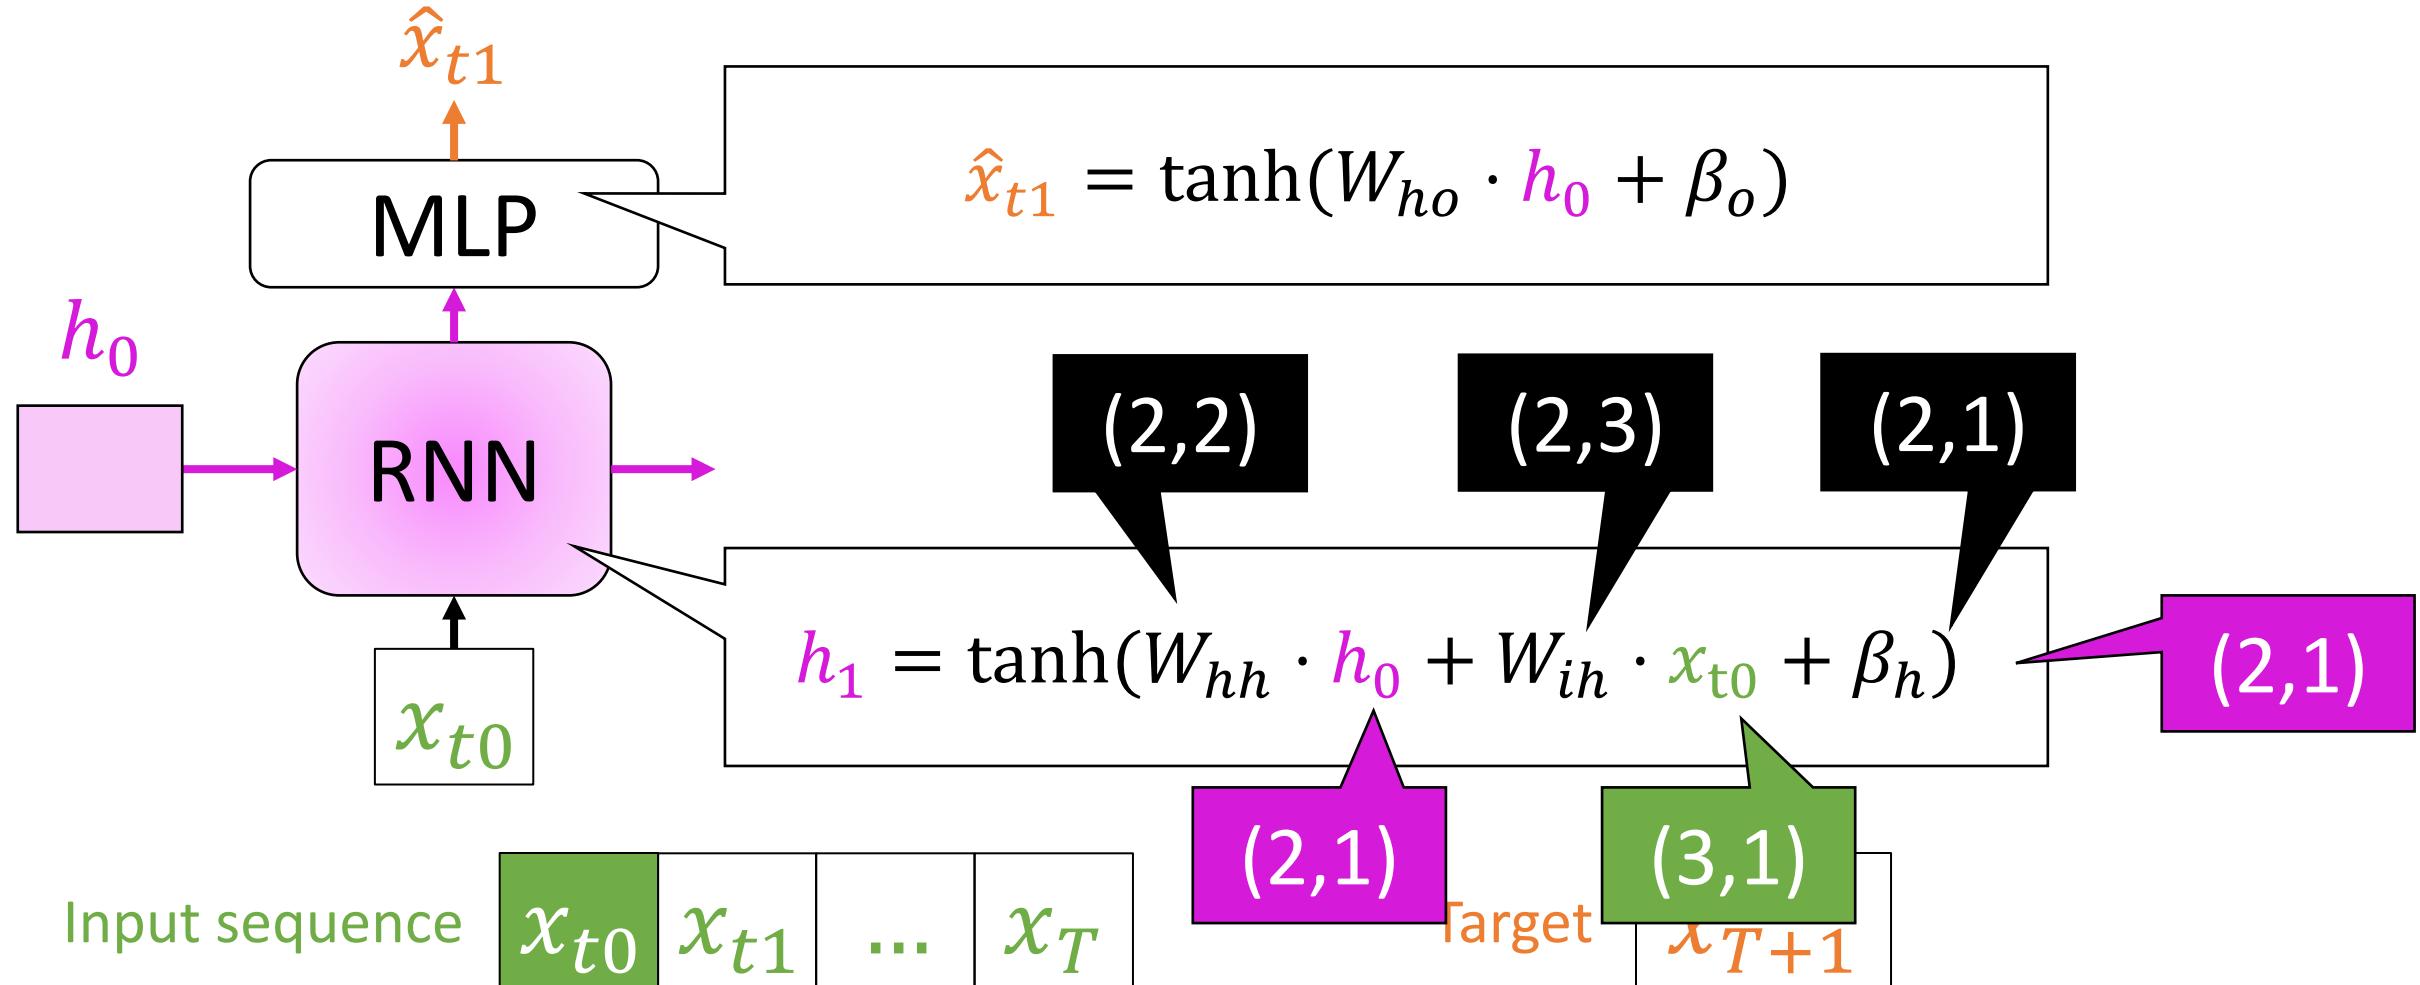
Target



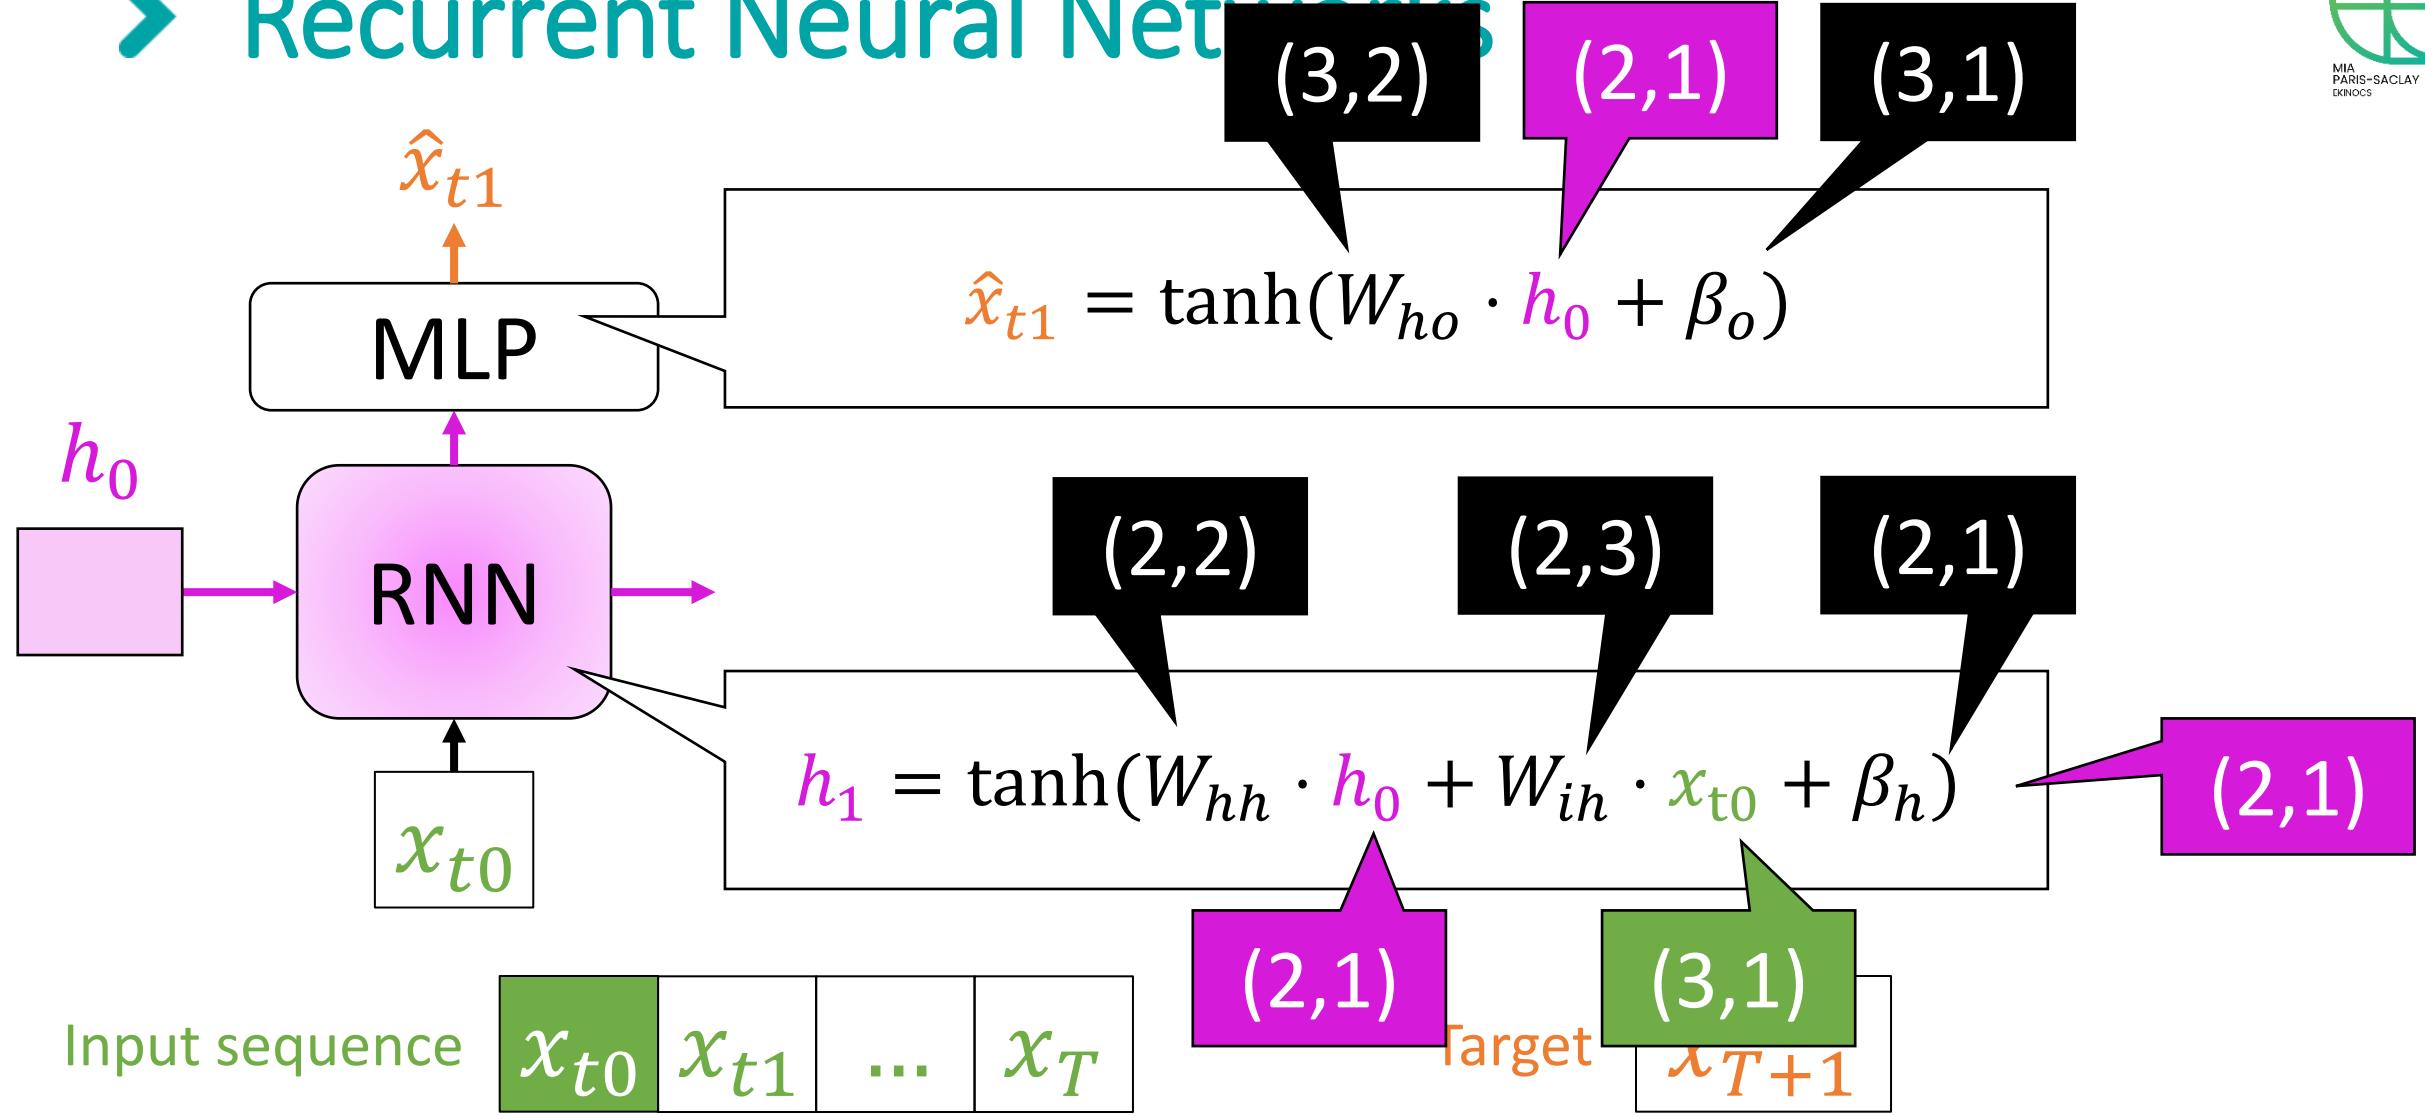
# > Recurrent Neural Networks



# > Recurrent Neural Networks



# > Recurrent Neural Networks



# > Issues with RNNs

- Unrolling is a simple and effective technique
  - However, it factually creates *very deep* networks
  - Backpropagation through 1,000 layers can be an issue
- Exploding gradient, e.g.  $(1.01)^{1000}$ 
  - Values so big that it's an issue representing them
  - But not only, HUGE gradient updates
- Vanishing gradient, e.g.  $(0.99)^{1000}$ 
  - Same issues as above with internal representation
  - And super-small gradient updates (run out of epochs)

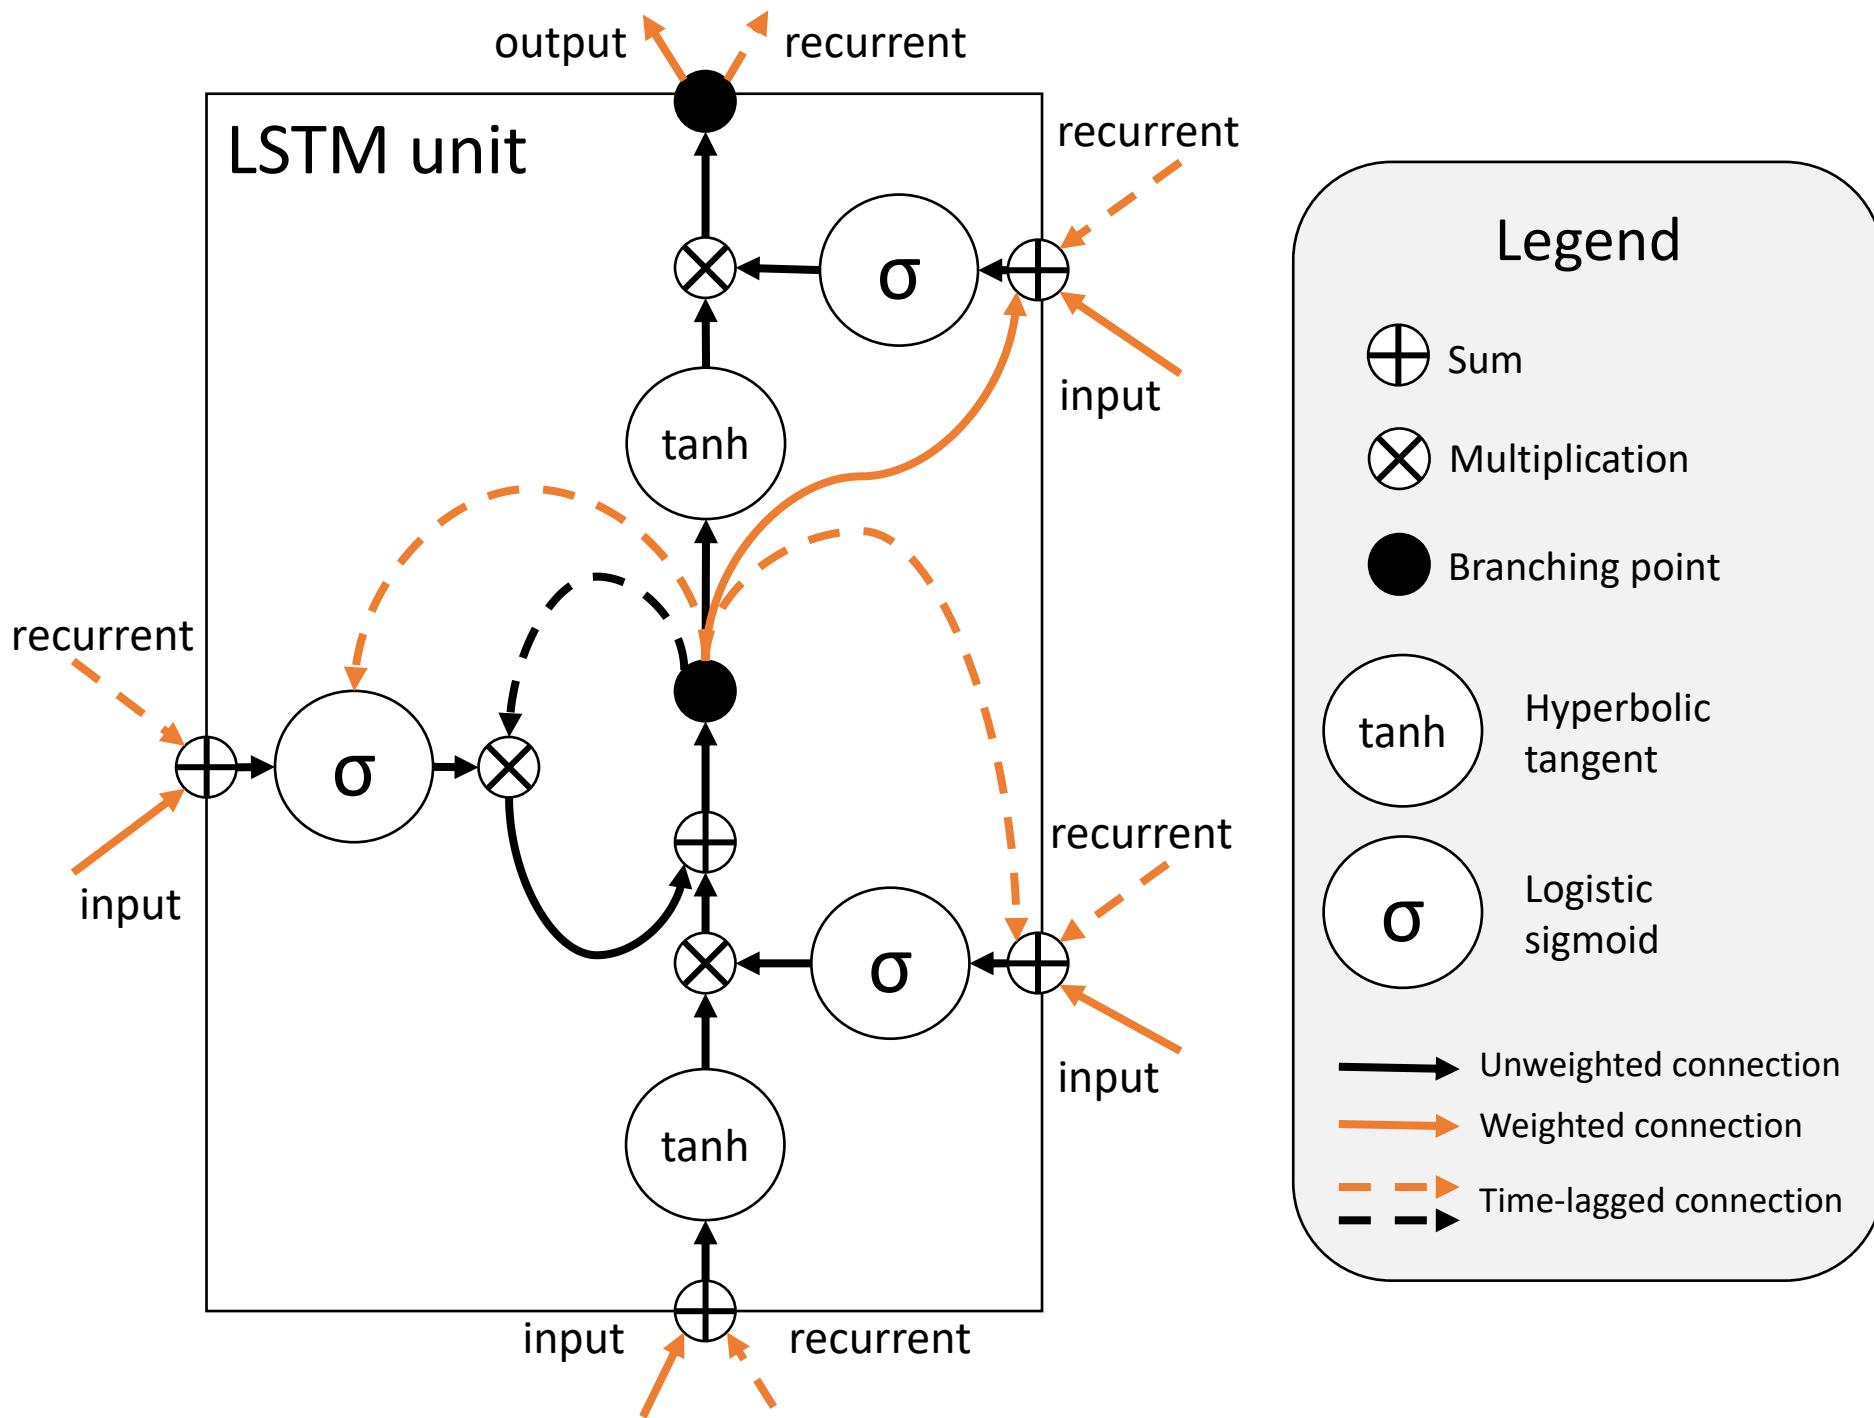
# > Issues with RNNs

- Solutions have been proposed
  - **Gradient clipping** to solve exploding gradient, normalize

$$\nabla_w E = \begin{cases} \frac{\nabla_w E}{\|\nabla_w E\|} & \text{if: } \|\nabla_w E\| > 1 \\ \nabla_w E & \text{else} \end{cases}$$

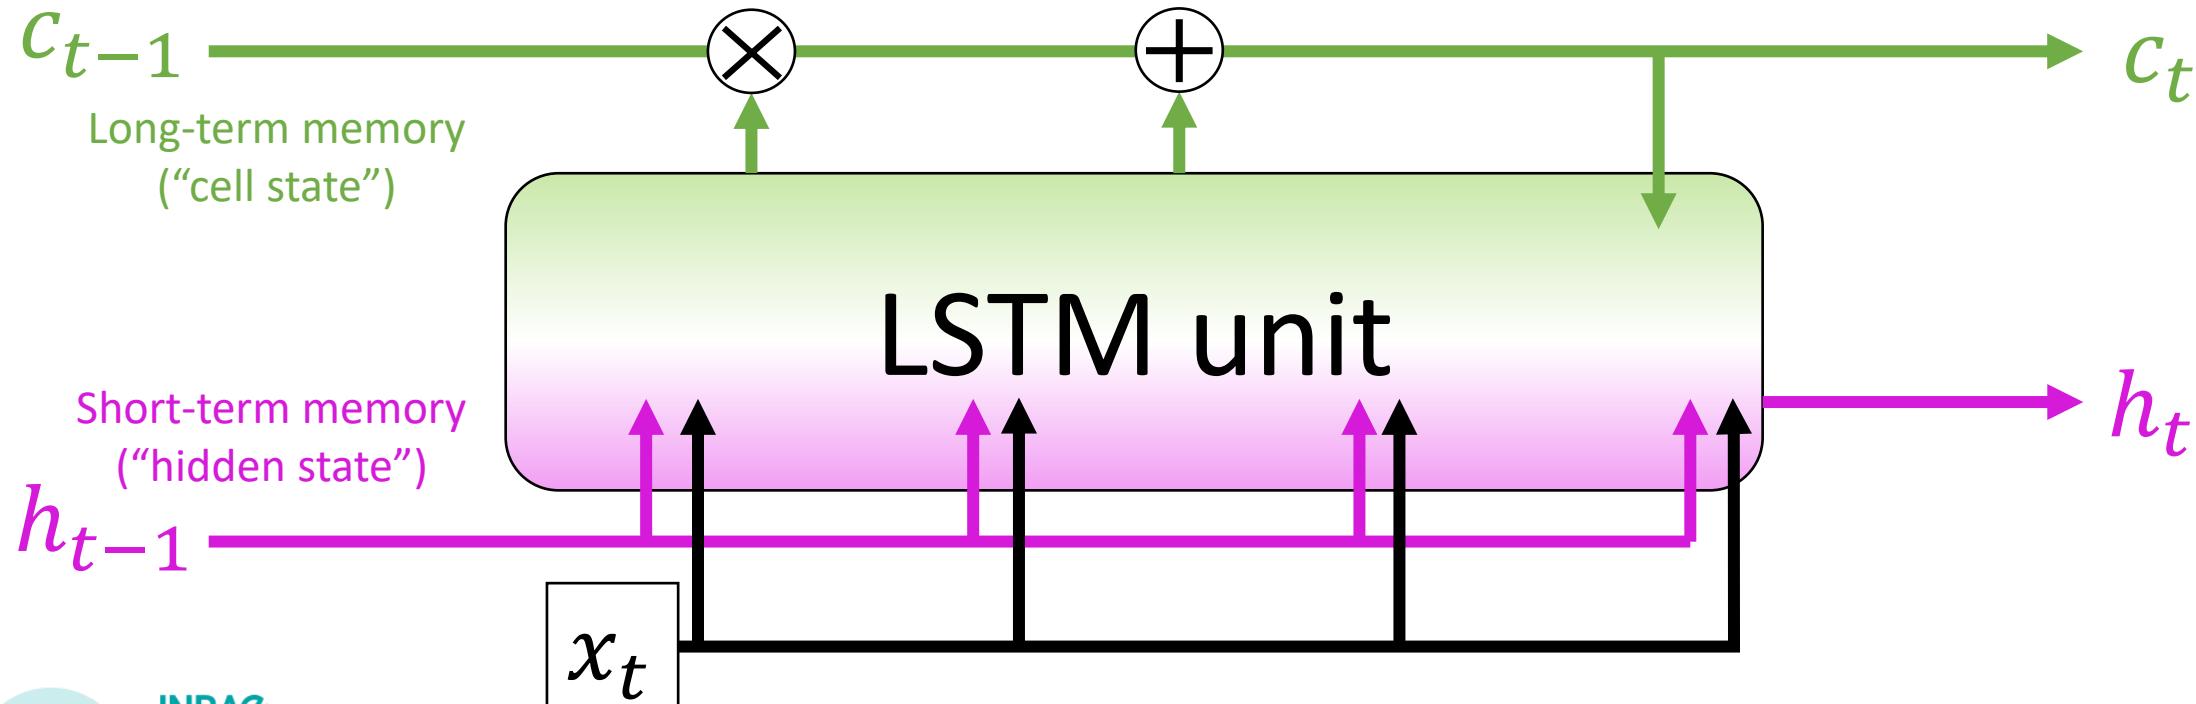
E here is the error function (loss)

- Long-Short Term Memory Networks (LSTMs)
  - New type of module, “reset” hidden state when needed
  - No weights on the path of updating the long-term history
  - Successful in practical applications

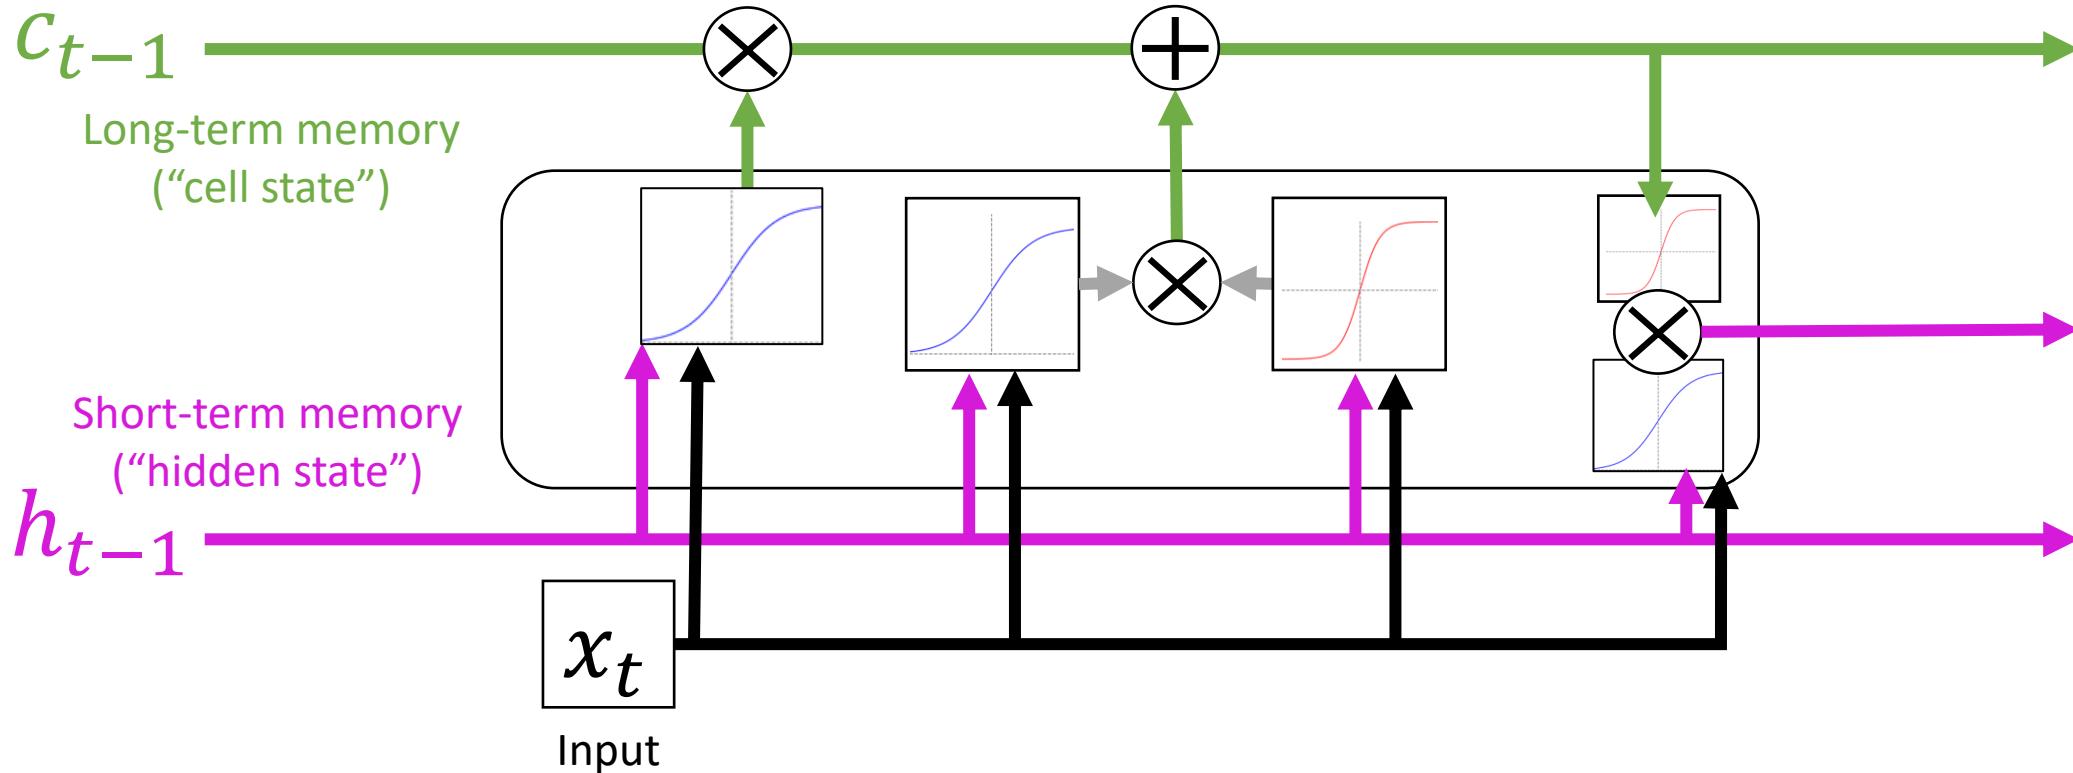


# > Long-Short Term Memory Networks

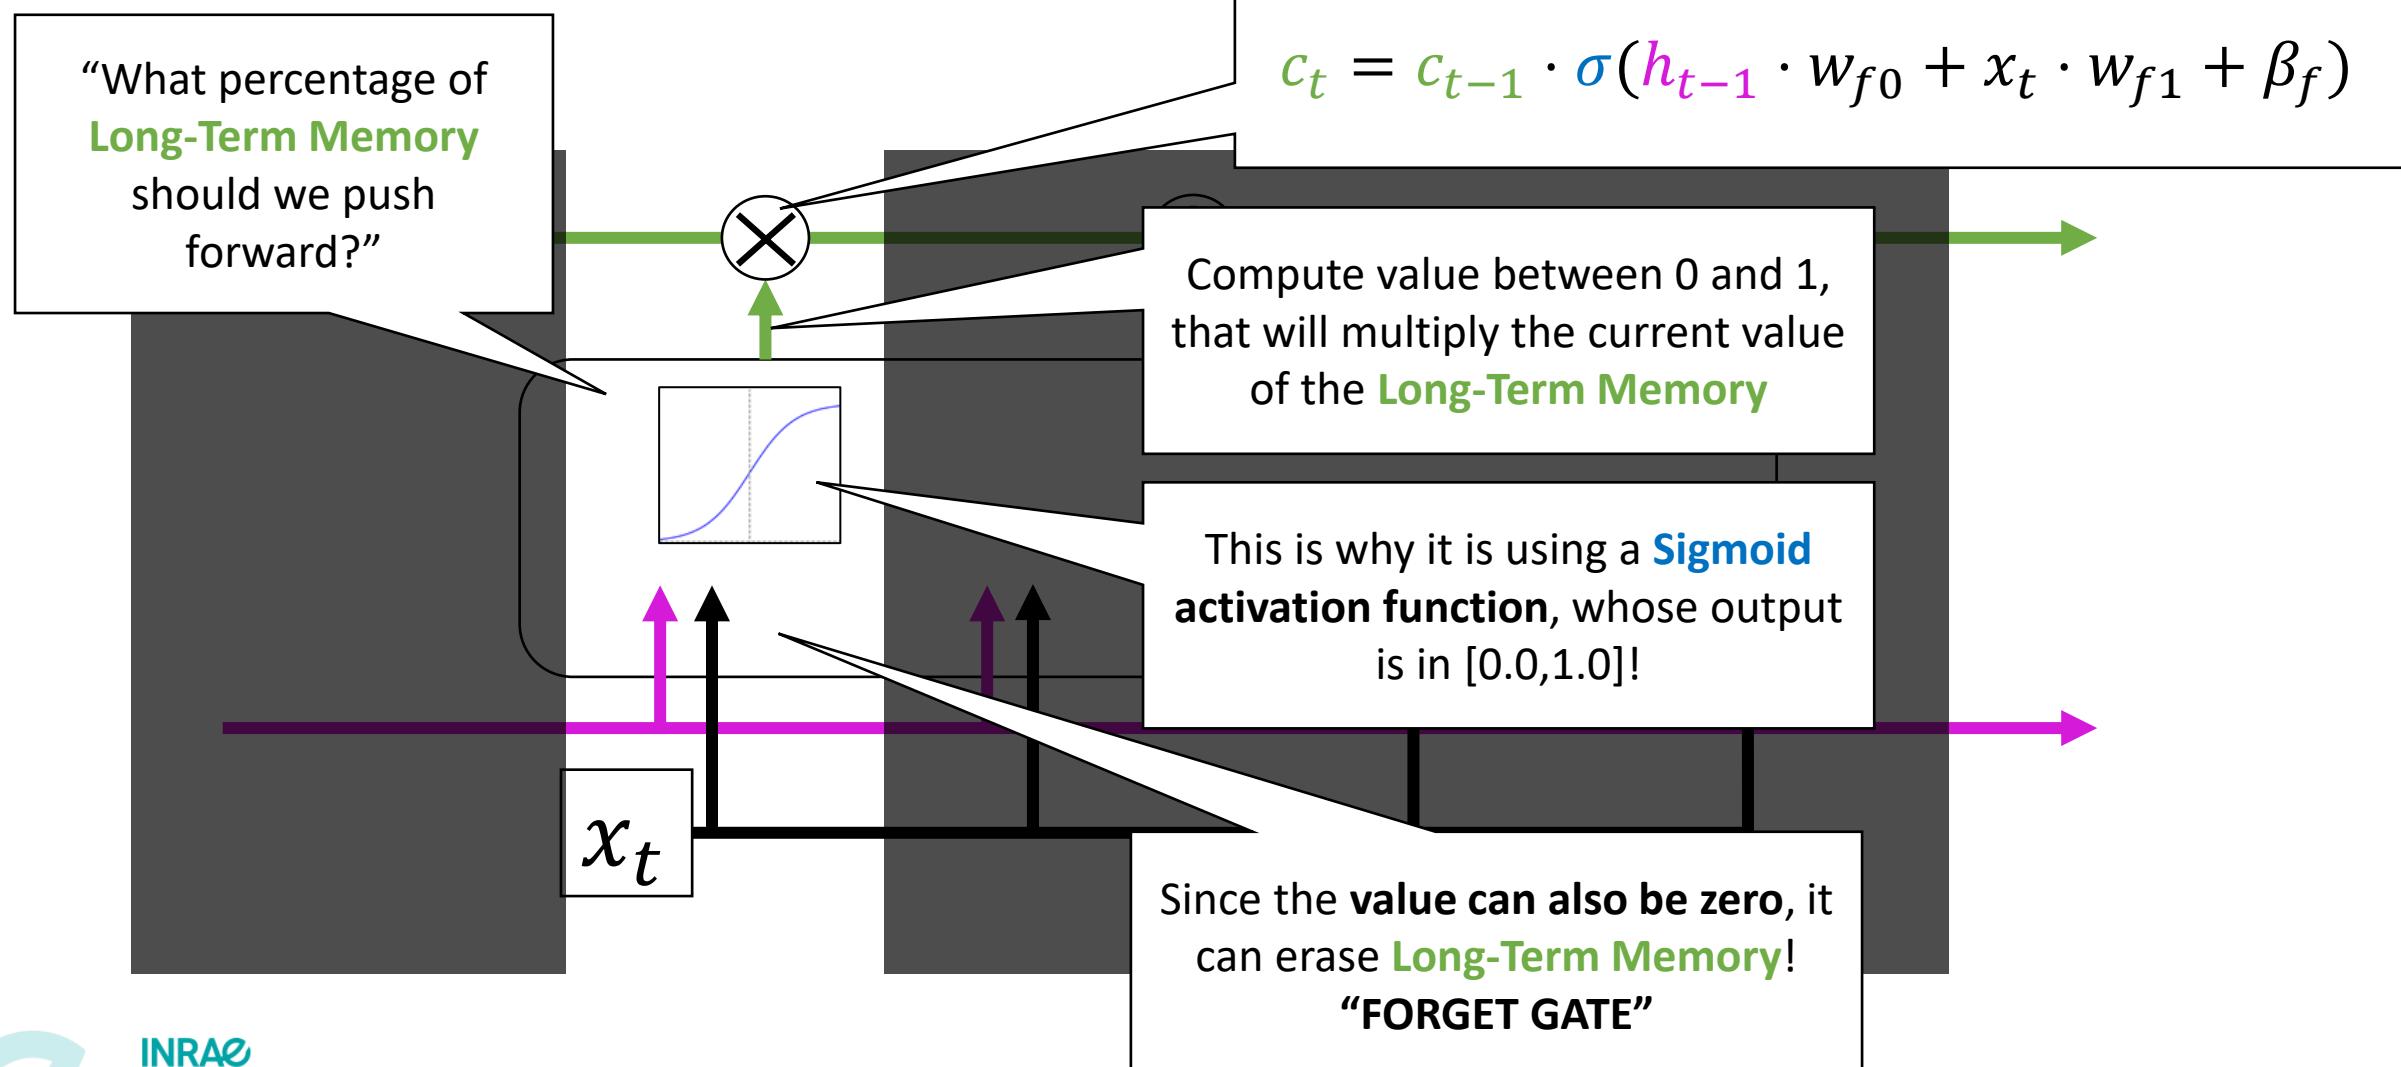
- Two “memory lanes” per unit
  - **Long-term memory**, no weights (!), avoids exploding/vanishing
  - **Short-term memory**, carried out from the recent past



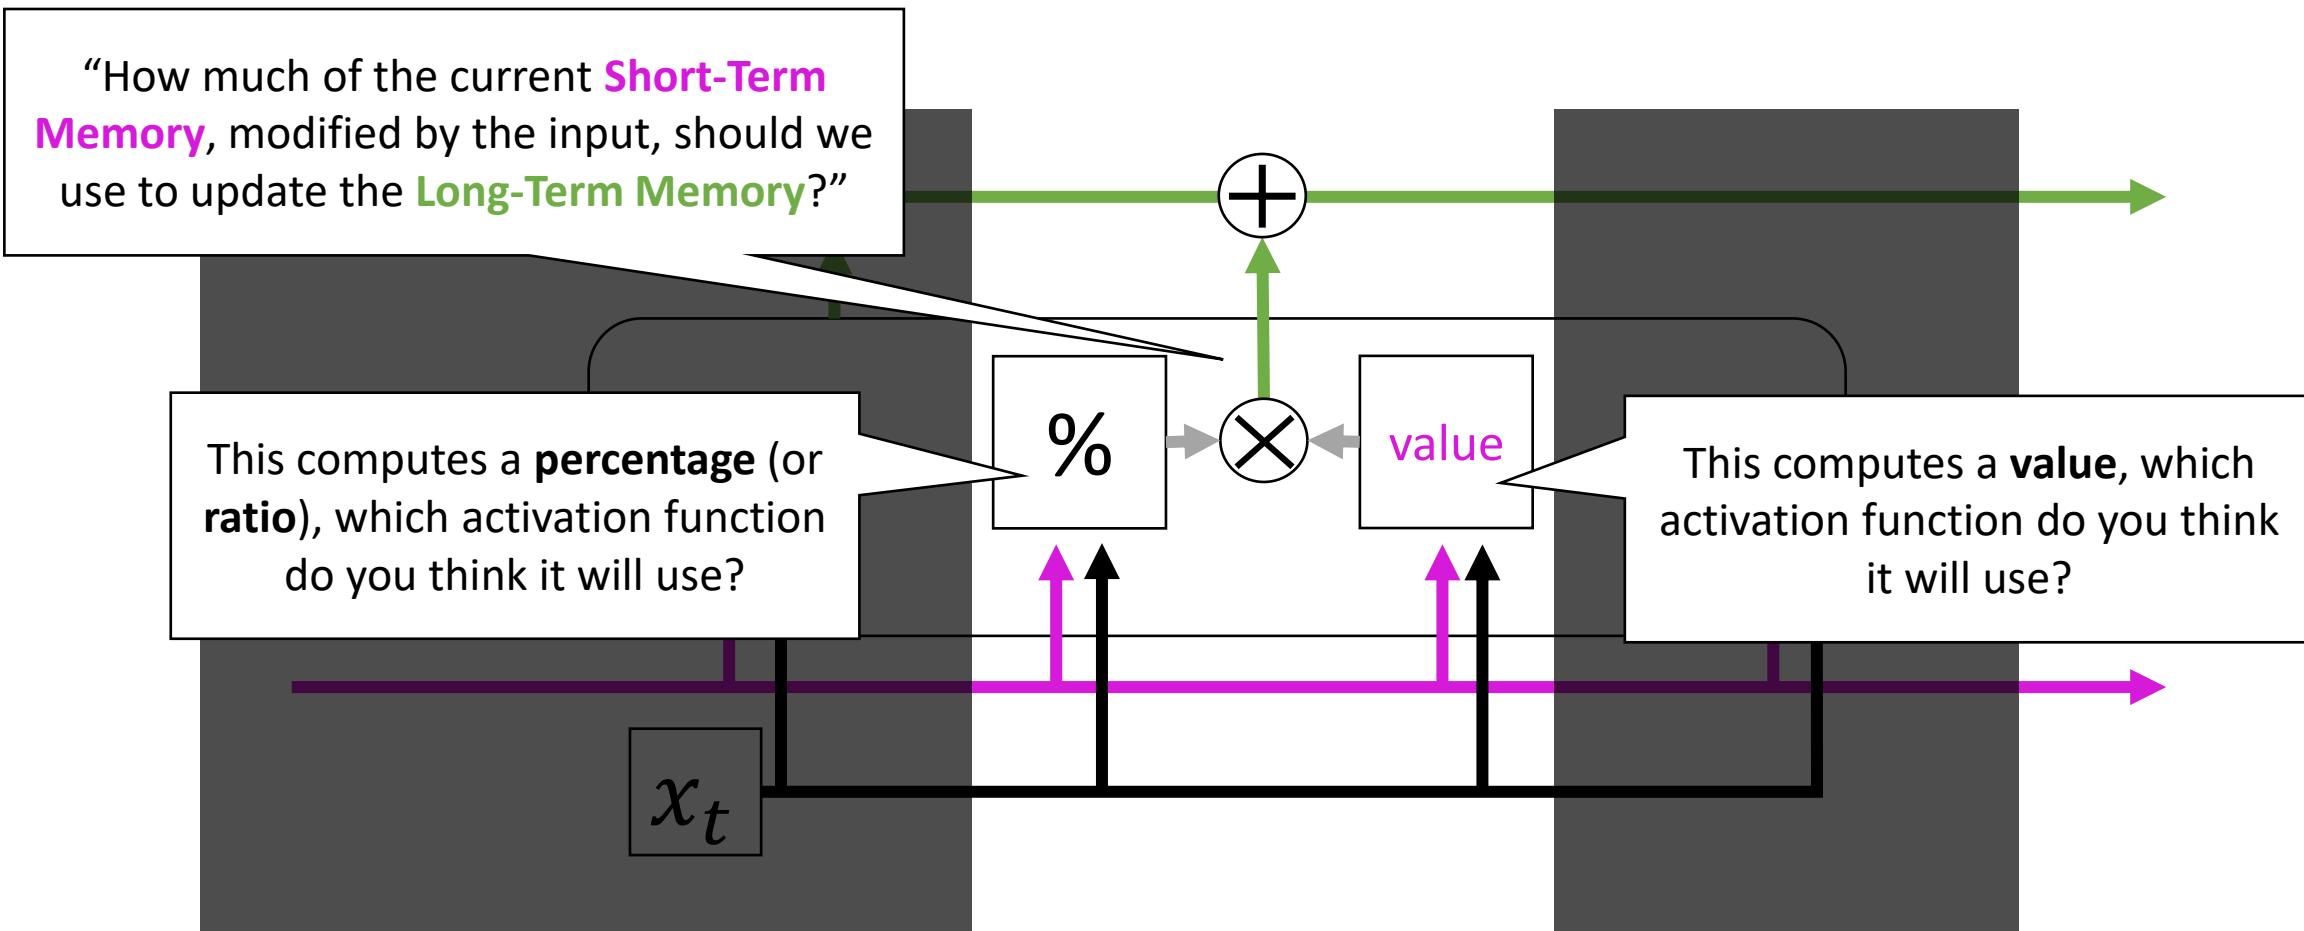
# > Long-Short Term Memory Unit



# > Long-Short Term Memory Unit



# > Long-Short Term Memory Unit



# > Long-Short Term Memory Unit

"How much of the current **Short-Term Memory**, modified by the input, should we use to update the **Long-Term Memory**?"

$$c_t = c_{t-1} + \sigma(h_{t-1} \cdot w_{i0} + x_t \cdot w_{i1} + \beta_{i1}) \cdot \tanh(h_{t-1} \cdot w_{i2} + x_t \cdot w_{i3} + \beta_{i2})$$

$$\sigma(h_{t-1} \cdot w_{i0} + x_t \cdot w_{i1} + \beta_{i1})$$

$$\tanh(h_{t-1} \cdot w_{i2} + x_t \cdot w_{i3} + \beta_{i2})$$

$x_t$

"INPUT GATE"

# > Parenthesis

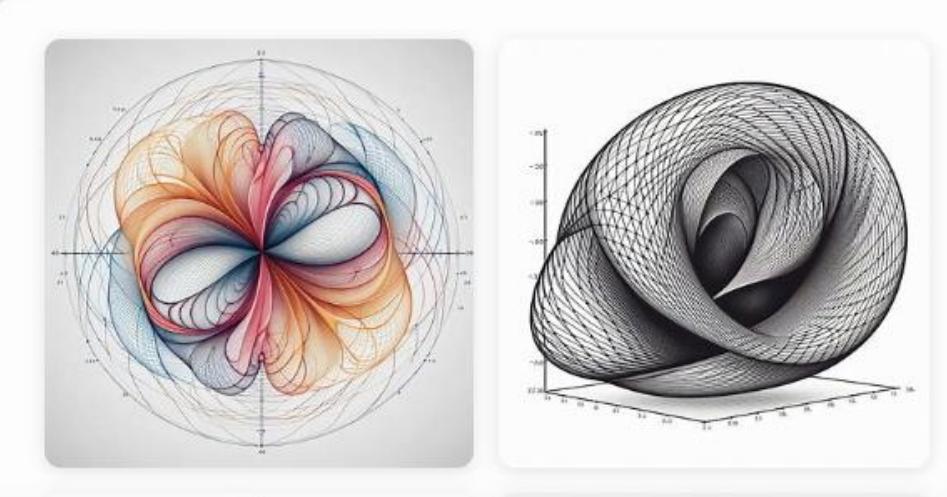
 Copilot

Certainly! Here's the hyperbolic tangent function ( $\tanh$ ) graph between -1 and 1, with unlabeled axes:

$$f(x) = \tanh(x)$$

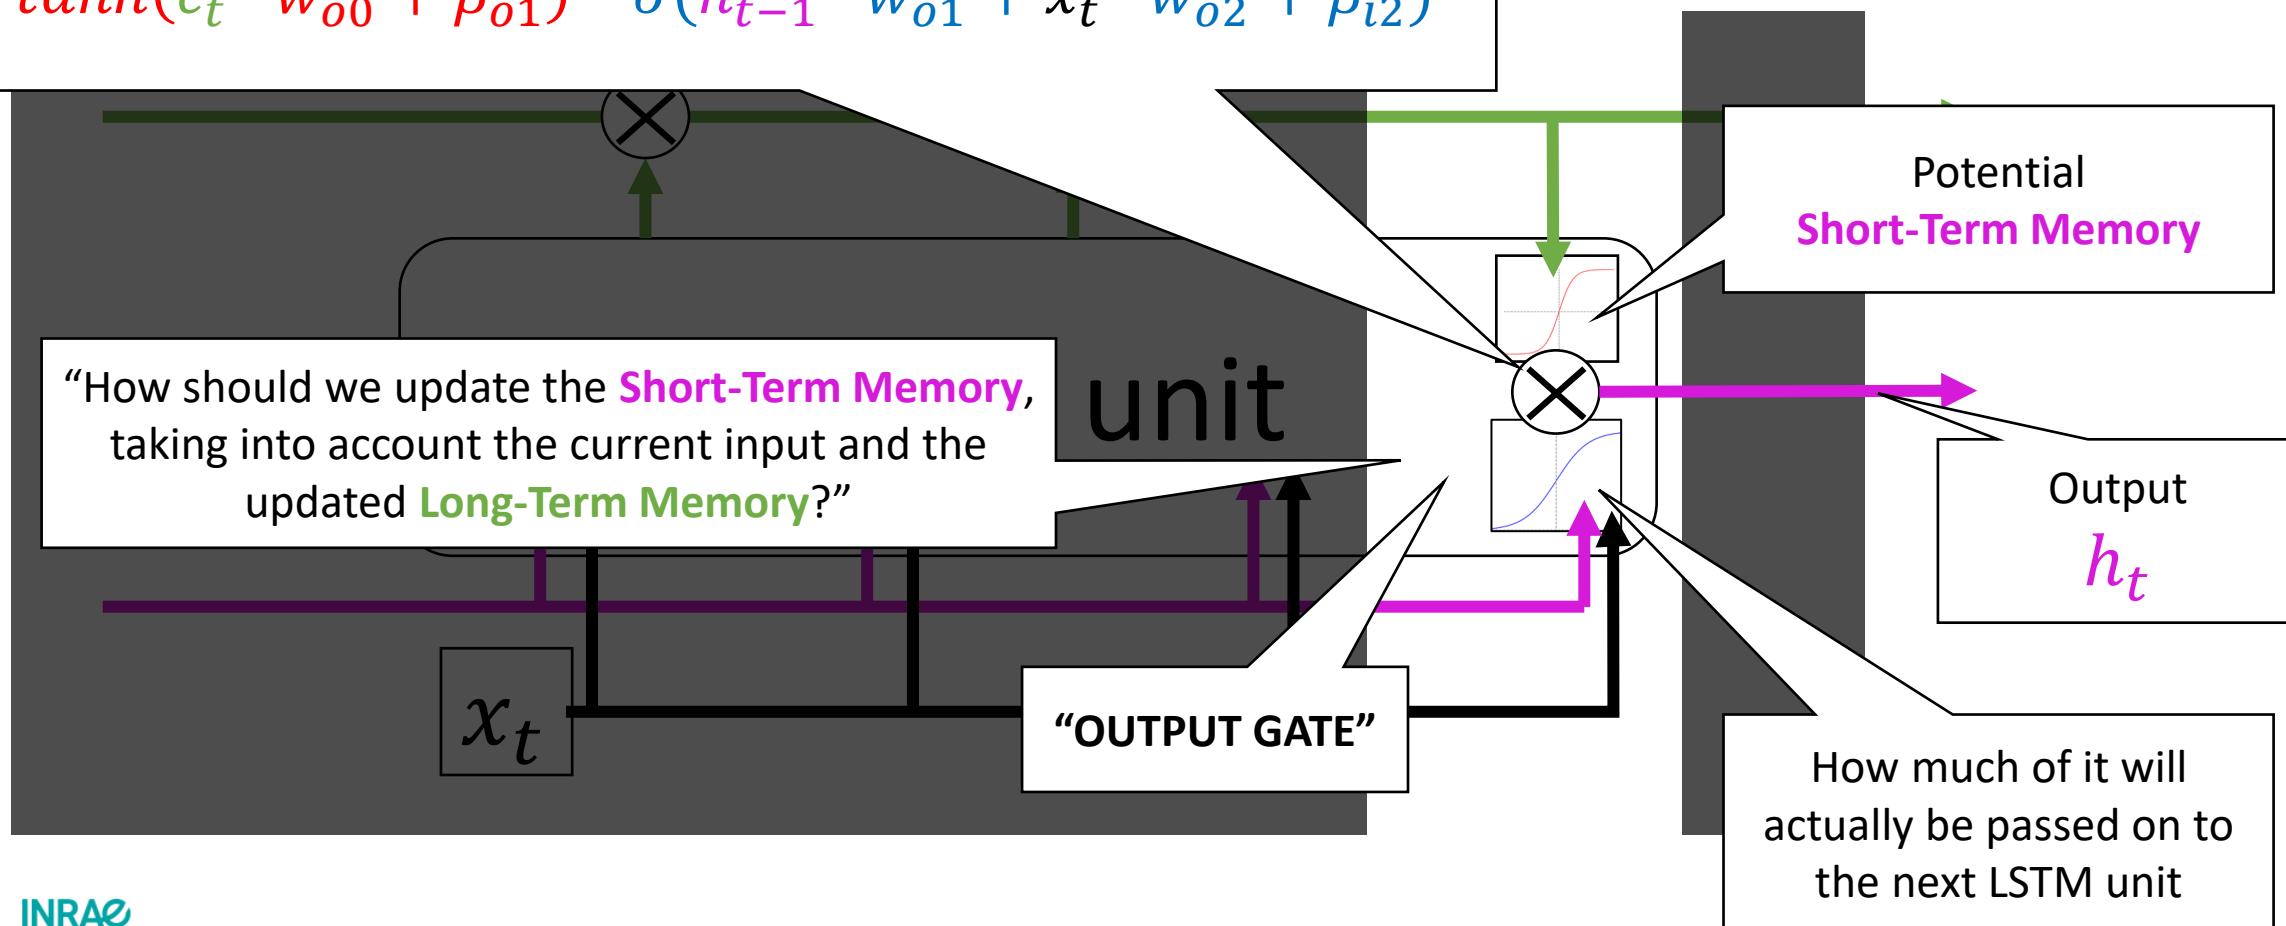
!Hyperbolic Tangent Graph

Like Dislike Share Download Sound 2 of 30

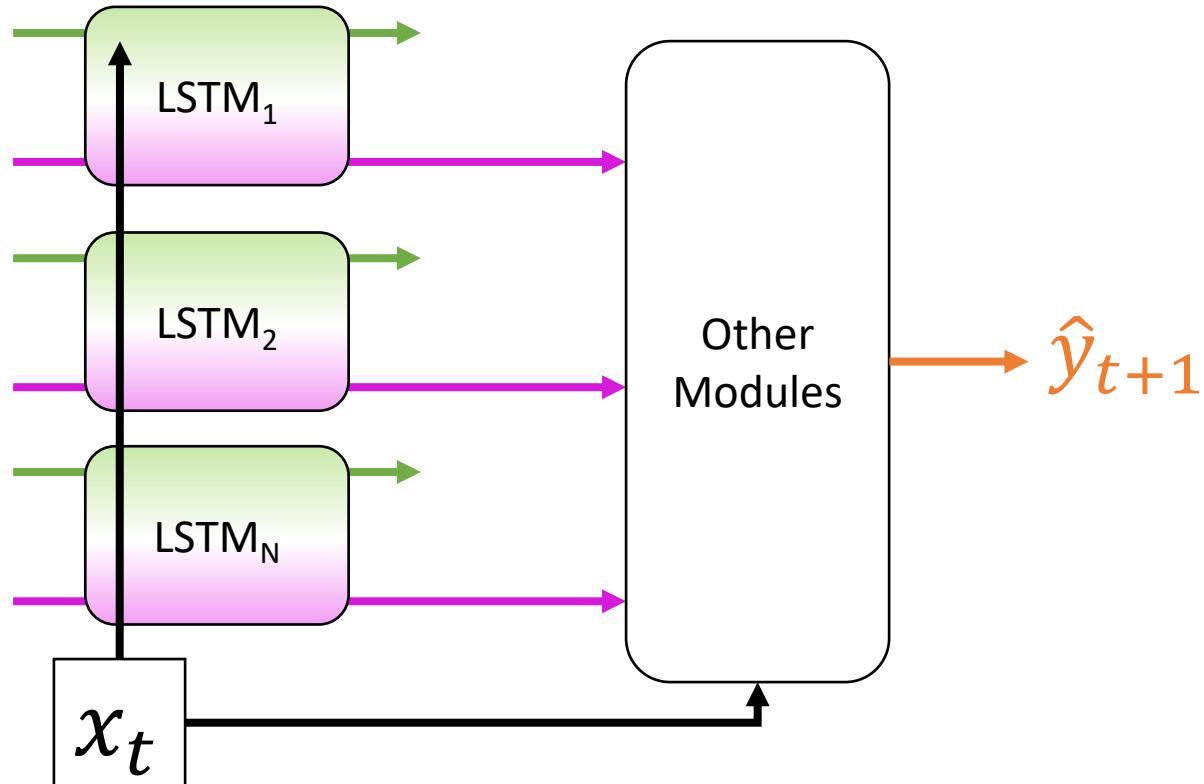


# > Long-Short Term Memory Unit

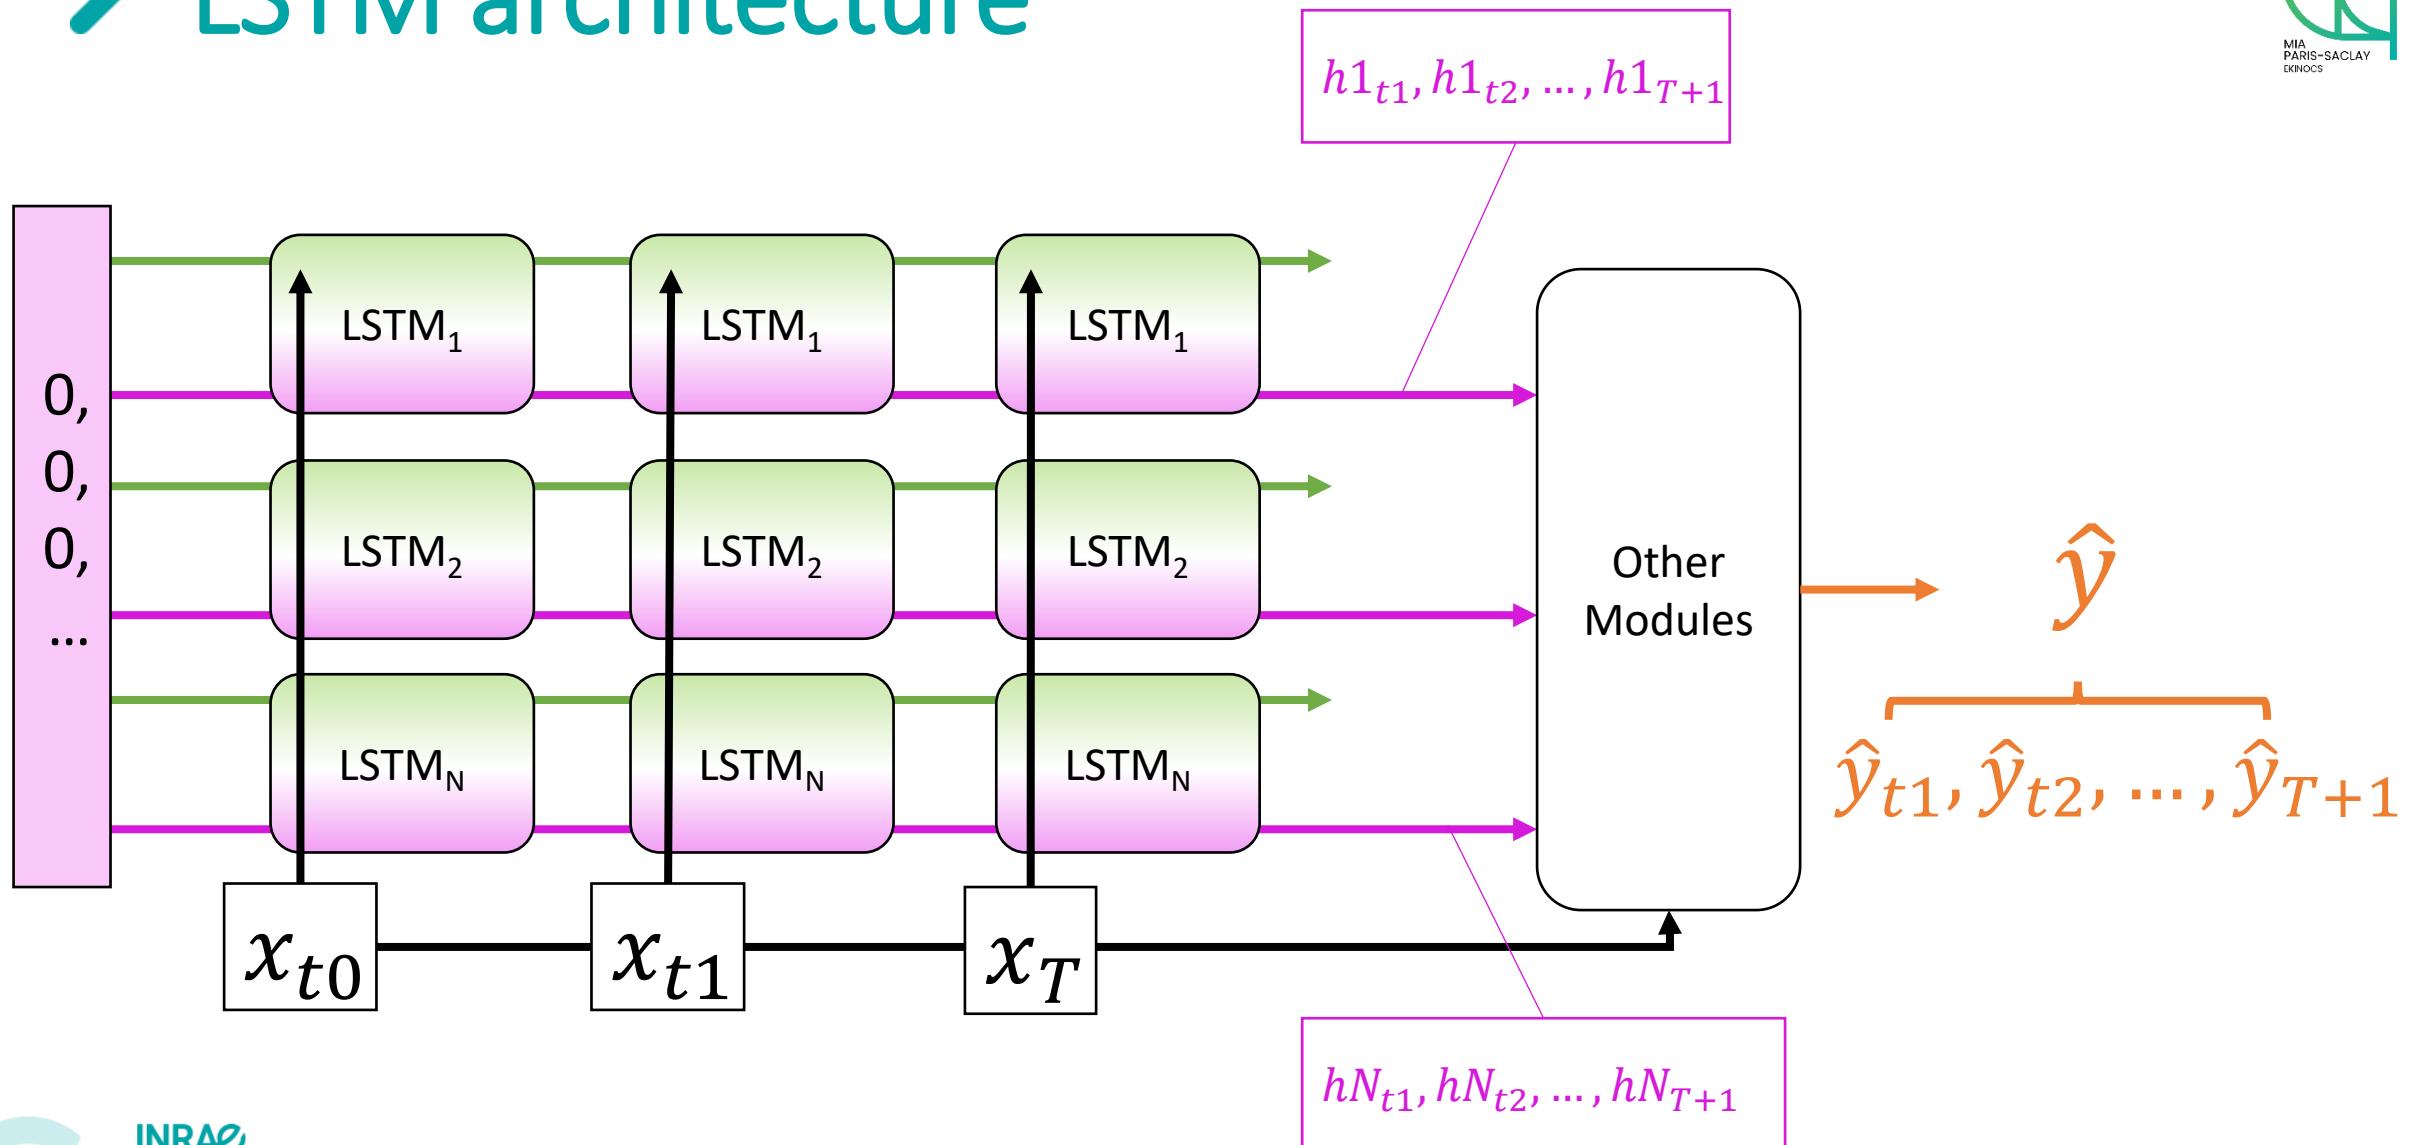
$$h_t = \tanh(c_t \cdot w_{o0} + \beta_{o1}) \cdot \sigma(h_{t-1} \cdot w_{o1} + x_t \cdot w_{o2} + \beta_{i2})$$



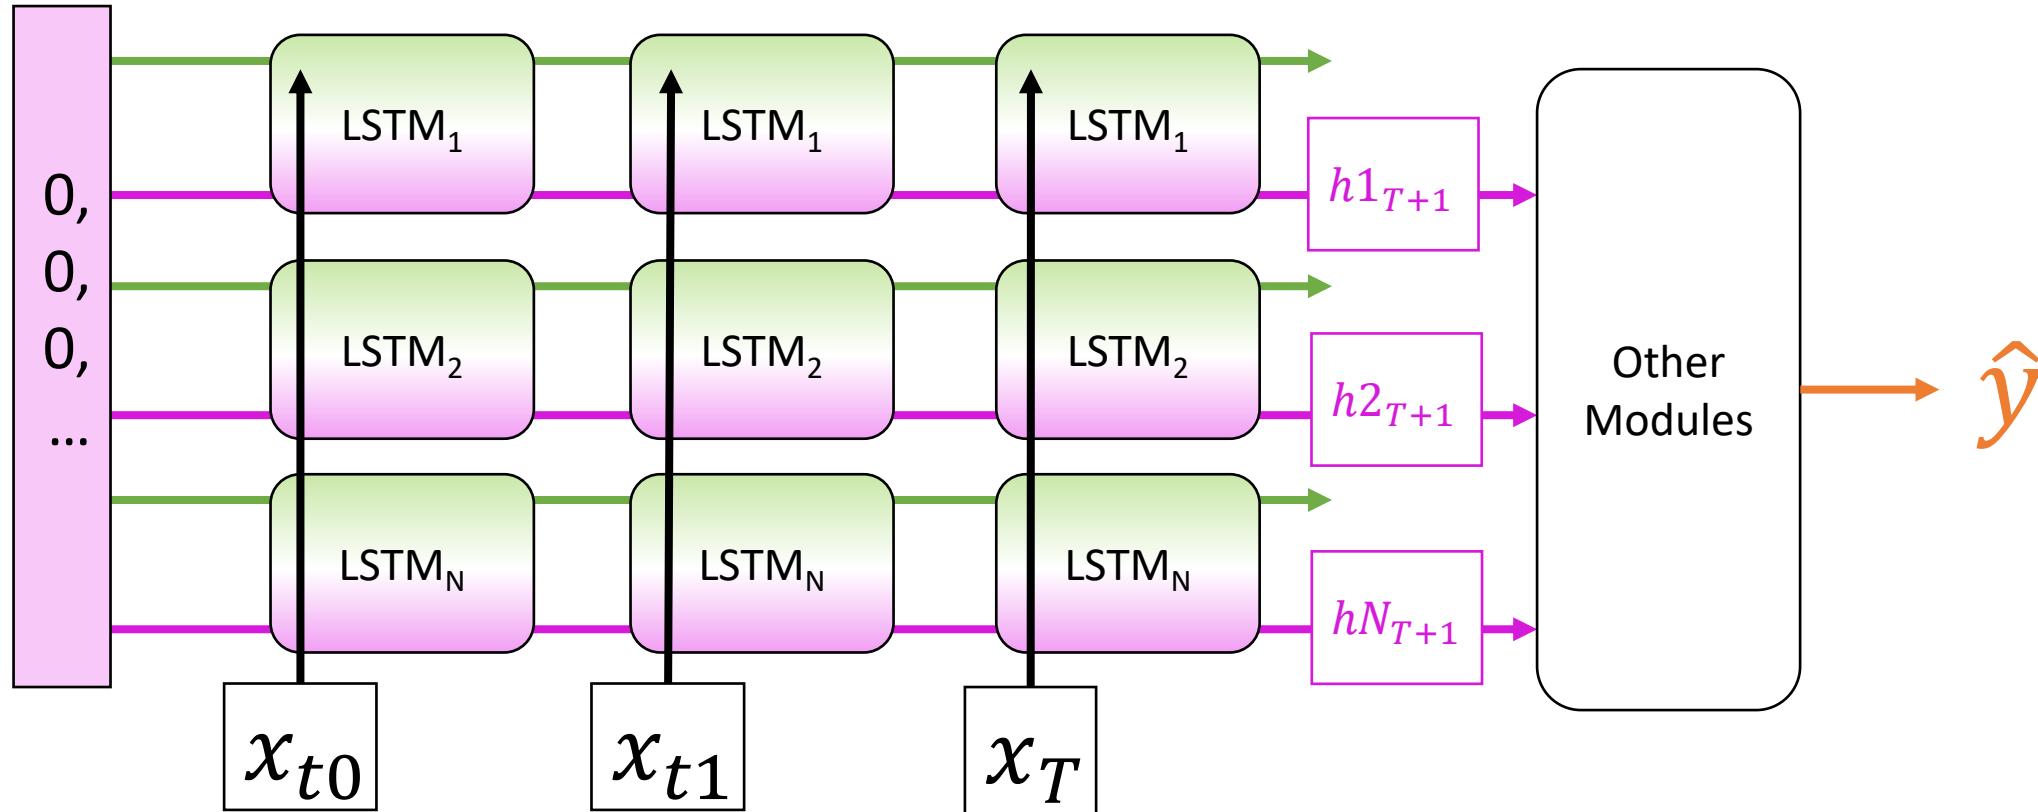
# > LSTM architecture



# > LSTM architecture



# > LSTM architecture



# > Gated Recurrent Units (GRUs)

- Variant of LSTM with less parameters per unit (12 vs 16)
  - Considered more or less just as accurate, slightly faster to train
  - More recent applications favor GRUs over LSTMs
  - But not always, in practice people *try both* and then pick

## pytorch RNN

$$h_t = \tanh(x_t W_{ih}^T + b_{ih} + h_{t-1} W_{hh}^T + b_{hh})$$

## pytorch LSTM

$$\begin{aligned} i_t &= \sigma(W_{ii}x_t + b_{ii} + W_{hi}h_{t-1} + b_{hi}) \\ f_t &= \sigma(W_{if}x_t + b_{if} + W_{hf}h_{t-1} + b_{hf}) \\ g_t &= \tanh(W_{ig}x_t + b_{ig} + W_{hg}h_{t-1} + b_{hg}) \\ o_t &= \sigma(W_{io}x_t + b_{io} + W_{ho}h_{t-1} + b_{ho}) \\ c_t &= f_t \odot c_{t-1} + i_t \odot g_t \\ h_t &= o_t \odot \tanh(c_t) \end{aligned}$$

## pytorch GRU

$$\begin{aligned} r_t &= \sigma(W_{ir}x_t + b_{ir} + W_{hr}h_{(t-1)} + b_{hr}) \\ z_t &= \sigma(W_{iz}x_t + b_{iz} + W_{hz}h_{(t-1)} + b_{hz}) \\ n_t &= \tanh(W_{in}x_t + b_{in} + r_t \odot (W_{hn}h_{(t-1)} + b_{hn})) \\ h_t &= (1 - z_t) \odot n_t + z_t \odot h_{(t-1)} \end{aligned}$$



# Questions?

## Bibliography

- Amari, S. I. (1972). *Learning patterns and pattern sequences by self-organizing nets of threshold elements.* IEEE Transactions on computers, 100(11), 1197-1206. [[RNNs](#)]
- Cho, K., Van Merriënboer, B., Bahdanau, D., & Bengio, Y. (2014). On the properties of neural machine translation: Encoder-decoder approaches. arXiv preprint arXiv:1409.1259. [[GRUs](#)]
- Hochreiter, S., & Schmidhuber, J. (1997). *Long short-term memory.* Neural computation, 9(8), 1735-1780. [[LSTMs](#)]
- Hopfield, J. J. (1982). Neural networks and physical systems with emergent collective computational abilities. Proceedings of the national academy of sciences, 79(8), 2554-2558. [[RNNs](#)]
- StatQuest, Long Short-Term Memory (LSTM), Clearly Explained,  
<https://www.youtube.com/watch?v=YCzL96nL7j0>

Images and videos: unless otherwise stated, I stole them from the Internet. I hope they are not copyrighted, or that their use falls under the Fair Use clause, and if not, I am sorry. Please don't sue me.