

INRAE



université
PARIS-SACLAY

➤ Neural networks and Deep learning

Alberto TONDA


*UMR 518 MIA-PS, INRAE, AgroParisTech, Université Paris-Saclay
UAR 3611, Institut des Systèmes Complexes de Paris Île-de-France*

➤ Outline

- Artificial neural networks
- Optimizing a neural network
- Overparametrization (“double descent”)
- Convolutional neural networks
- Recurrent neural networks
- Transformers
- Autoencoders
- From the point of view of optimization...




➤ Artificial neural networks




➤ Artificial neural networks


- Based on an old (and wrong) model of a real neuron



➤ Artificial neural networks



➤ Artificial neural networks



➤ Artificial neural network

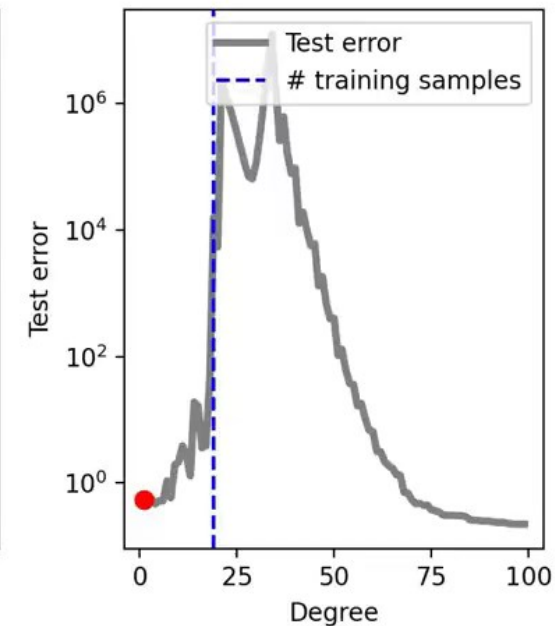
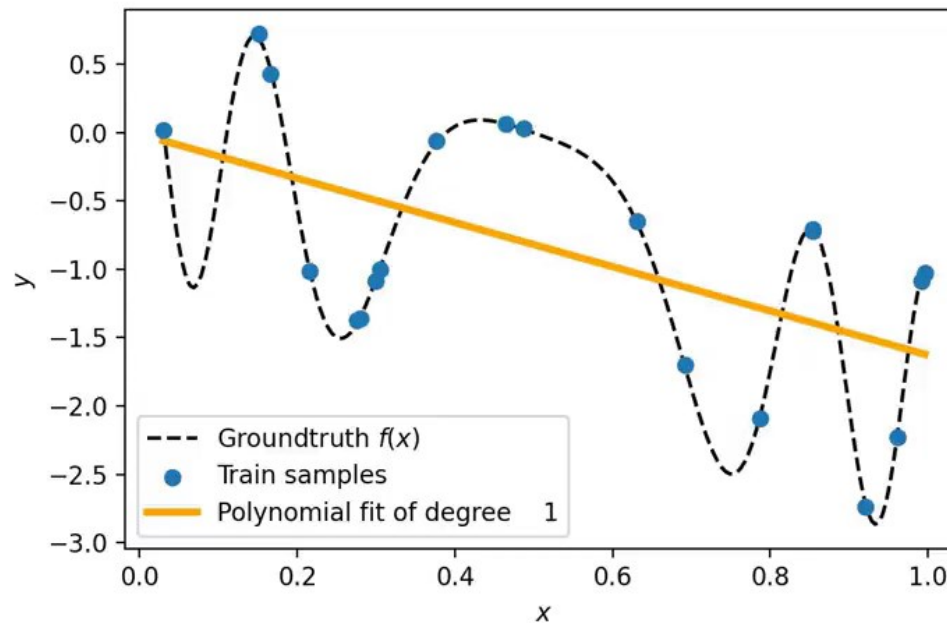
- So, an ANN is just a *very complex function*
- In ML terms, its parameters are the weights
- Modern ANNs have thousands/millions/**billions** of weights!
- How to optimize a function in such high dimension?

➤ SGD and backpropagation

- Objective function is called **loss function**
 - To be **minimized**; exact content depends on the task
 - MSE for regression, categorical cross-entropy for classification
- It is possible to compute the derivative of the loss
 - Thanks to chain rule and backpropagation
 - Partial derivative of loss with respect to **each weight/parameter**
 - Stochastic Gradient Descent (and successors)
 - “Stochastic” = use only a subset of samples at each iteration


➤ Overparametrization

- Wait a second!
 - From ML, we learned that having too many parameters is **bad**!
 - Models with too many parameters tend to **overfit** terribly...right?



➤ Overparametrization

- What is happening here? Well, we don't really know
 - Empirical results, **overparametrizing** improves **generalization**
 - “Double descent” or “W figure”
 - Does **regularization** play a role?




➤ Why are neural networks so successful?

- Extremely performant for **relational data**
 - Thanks to recent architectures (CNNs, RNNs, Transformers, ...)
 - Images, text, sound, ...
 - Other ML approaches are just not as good
- Graphics Processing Units (GPUs)
 - Are really good at performing parallel computations
 - And in fact, they are excellent to speed up gradient computation
- Multiple outputs can be interpreted as pixels/audio/text...
- More and more data is available!



➤ Convolutional neural networks

- Specialized layers that scan the whole structured input
- They can find patterns in every part of it
- Example: images (square window, slide over pixels)

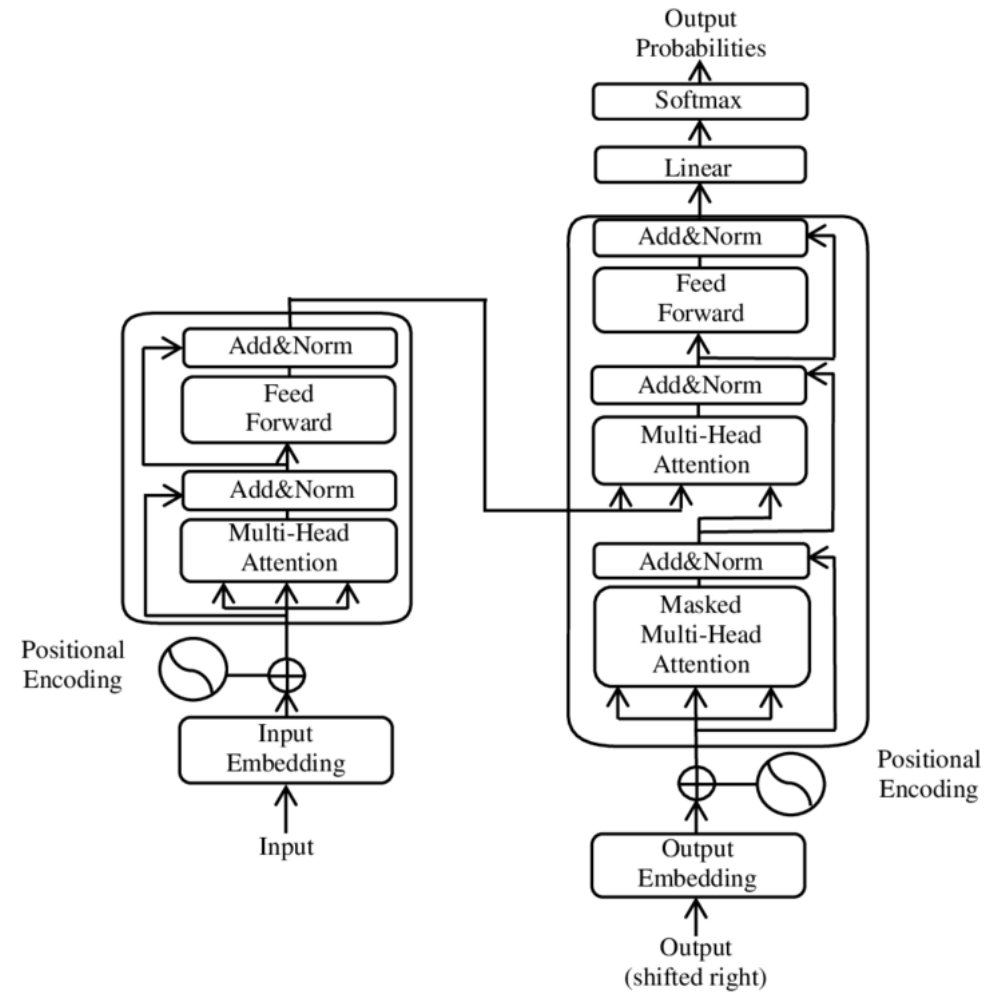


➤ Recurrent neural networks

- Current output is not just $y_n = f(\mathbf{x}_n)$
 - Instead, there is a **state** of the system $y_n = f(\mathbf{x}_n, S)$
 - $S = g(\mathbf{x}_0, \mathbf{x}_1, \dots, \mathbf{x}_{n-1})$
 - Sequences, time series, dynamic systems...
- RNNs use specialized layers to **keep a memory** of the state
- Modern developments
 - Long-Short Memory Networks (LSTMs)
 - Gated Recurrent Units (GRUs)

➤ Transformers

- De-facto replaced RNNs
- Gets all the data at once
- Specialized layers that try to capture relationships between parts of the input



➤ Autoencoders


- Train unsupervised to **exactly reproduce the input**

Optimization task: $\operatorname{argmin}(\sum_{i=0}^N | \hat{x}(i) - x(i) |)$



➤ Autoencoders


- Force the computation to go through bottleneck
- Dimension of the bottleneck much smaller than input



Bottleneck/Latent space


➤ Autoencoders

- Remove the second part of the model




➤ Autoencoders

- Remove the second part of the model
- Going from input to (lower dimensionality) bottleneck




➤ Autoencoders

- However, we can also remove the first part
- From the bottleneck/latent space to the output




➤ Autoencoders


- However, we can also remove the first part
- From the bottleneck/latent space to the output




➤ Autoencoders and generative NNs



➤ Autoencoders and generative NNs



➤ Autoencoders and generative NNs



➤ From the point of view of optimization...

- Interesting results from a 2016 paper, “Understanding deep learning requires rethinking generalization”
 - Shows that a neural network can *memorize* a dataset
 - In other words, it has enough **capacity** to overfit completely
 - But when there is an actual relationship $y = f(X)$, **finds it**
- From an optimization point of view
 - There is a **global optimum** of the weights value
 - Global optimum corresponds to **total overfit**
 - SGD finds a *local optimum* that has **better generalization (!)**

INRAE



université
PARIS-SACLAY

➤ Questions?

Bibliography

- Goodfellow et al., *The Deep Learning Book*, 2016
- Zhang et al., *Understanding Deep Learning Requires Rethinking Generalization*, 2016

Images and videos: unless otherwise stated, I stole them from the Internet. I hope they are not copyrighted, or that their use falls under the Fair Use clause, and if not, I am sorry. Please don't sue me.